



NURSE AI OS • SETUP GUIDE 2 OF 2

The Hermes Power Guide

For power users and nurse leaders: run a unit of isolated agents, schedule work that does itself, route models like an architect — with EDENA governance embedded in every layer.

Agents propose. Humans judge. Nurses steward.

Free for nurses and nursing students • Founding Year 2026 • nurse-ai-os.org

Who this guide is for

You have the basics running — SOUL file loaded, Starter Kit in place, first dashboard built — and you are ready for the charge-nurse view: multiple agents with separate assignments, work that runs on a schedule without you, and the governance discipline that lets you (and your organization) trust all of it. If you have not done the basics, start with the Hermes Essentials guide and the Pre-Procedure Checklist first.

The architecture you are standing on

Power comes from the stack, not the model. Every capability in this guide runs inside the same governed architecture — with **EDENA (Ethical Design & Enablement for Nursing Augmentation)** embedded in the background as the conscience layer:

You — the steward — judgment, accountability, the final word
SOUL files — standing instructions per agent: values, boundaries, red lines
EDENA — the conscience layer — every automated action classified before it runs; humans remain responsible for judgment
Profiles & structure — isolated agents, governed folders, memory in its place
Hermes engine + models — open-source horsepower, models routed by task weight

-  **GREEN** — low risk, may proceed
-  **YELLOW** — conditional, you approve first
-  **RED** — human authorization mandatory, always

The rule that scales: the more you automate, the more the classification matters. A scheduled job that drafts your morning brief is Green. A job that sends anything to another human is Yellow — it drafts, you approve. Anything clinical, patient-adjacent, or license-touching is Red: it does not get automated, ever. Governance before autonomy.

Profiles — one machine, a unit of agents

A profile is a **separate Hermes home**: its own config, keys, SOUL file, memory, sessions, skills, and scheduled jobs. Run a research agent, a writing agent, and a family-logistics agent on one machine without them mixing state — each with its own badge and its own med cart.

1 Create one

```
hermes profile create research # blank profile hermes profile create work --clone #
copy config, keys, SOUL, skills from current
```

Creating a profile also drops a shortcut, so **research chat** and **research setup** work directly.

2 Use and inspect it

```
hermes profile list # active profile marked with * hermes -p research chat # one-
off, doesn't change your default hermes profile use research # make it the sticky
default hermes profile show research # path, model, gateway status, skills
```

3 Set its model and working directory

```
research config set model.default anthropic/claude-sonnet-4 research config set
terminal.cwd /path/to/its/project research config set terminal.home_mode profile #
strict per-profile tool isolation
```

Give each profile its own SOUL.md — a research agent's boundaries differ from a family agent's.

Two things power users get wrong: (1) A profile is *not* a sandbox — it isolates Hermes state, not filesystem access; set **terminal.home_mode: profile** to tighten tool isolation. (2) Profile does not equal working directory — creating one does not set where its tools run; set **terminal.cwd** yourself. Official docs: hermes-agent.nousresearch.com/docs/user-guide/profiles

Cron — delegation that runs itself

A trigger, a prompt, a schedule. The running gateway checks for due jobs every 60 seconds, runs each in a fresh agent session, and saves output under `~/.hermes/cron/output/`. Jobs only fire while the gateway is running.

1

Create a morning briefing

```
hermes cron create "0 6 * * *" \ "Scan today's AI and healthcare-workforce
headlines. Keep the 3 that matter for a nurse leading AI adoption, summarize each in
2 lines, and save me the brief."
```

0 6 * * * means every day at 06:00. Other formats: one-shot delays (**30m**, **2h**), intervals (**every 6h**), full cron (**0 */6 * * ***).

2

Manage the schedule

```
hermes cron list # all jobs + next run hermes cron run <id> # run now to test hermes
cron pause <id> # pause / resume hermes cron edit <id> --schedule "0 6,18 * * *"
hermes cron remove <id>
```

!

Pair every job with a profile — and a tier

Run scheduled jobs under a dedicated profile so they act as their own isolated agent. Then classify the job before you schedule it: Green runs free; Yellow drafts and waits for you; Red does not get scheduled. Official docs: hermes-agent.nousresearch.com/docs/user-guide/features/cron

Route models like an architect

The architect does not lay the bricks. Put expensive reasoning where judgment lives, and cheap capacity where the work is routine:

Work type	Route to
Planning, strategy, judgment calls, review	Highest-capability reasoning model
Writing or building against a clear spec	Mid-tier model
Search, extraction, classification, summaries	Smallest model that reliably does it

Install the free **Architect Mode skill** (nurse-ai-os.org/setup-kit.html) and the discipline runs itself: five gates — Scope, Evidence, Attack, Verify, Report — on every high-stakes task, with confidence levels in every answer.

Governance habits for leaders

Keep a Kardex

End sessions with *"Give me a Kardex entry for this session"* — a five-line handoff note saved to memory. Automated work should leave the same paper trail your unit would expect.

Make corrections permanent

"Make that correction permanent" turns feedback into standing rules. Your agents get safer with use — and the rules travel in files you own.

Pilot, then scale

One agent, one workflow, two weeks, then review the outputs like an auditor before adding the next. The calm, defensible pace beats the impressive demo.

Taking this to a team or organization?

The same discipline scales: mapped burden, human gates, review rituals, and a one-month starter build live at **domondon-dominium.org**. Room & Nurse Assignment and Vacation Request programs — governed and human-gated — are coming for managers; early requesters shape the design.

The red lines do not move as you advance: no PHI, no patient-specific clinical decisions, no certification or deployment claims without human review and institutional governance. The more powerful your setup, the more these matter. Agents propose. Humans judge. Nurses steward.

Nurse AI OS · Founding Year 2026 · nurse-ai-os.org · Not medical advice · No PHI