

AI Operating System with the Heart of a Nurse

Starter Kit + Course Workbook

“She was the lady with the lamp. She was also the first to chart the truth in a way the powerful could not ignore — and the truth saved lives.”

Carry the lamp into the AI age

Florence Nightingale brought a lamp to the Crimean wards — and a notebook. She stayed present, comforted the dying, and protected human dignity at the bedside. She also gathered data, built her famous rose diagram, and proved that poor sanitation, not battle, was killing most soldiers. Compassion told her to act; **evidence told everyone else they had to.**

That union — **the caring heart and the counting mind** — is the whole point of this course. Data science and AI are not here to replace the nurse who notices. They are here to extend her reach, lighten her load, and let her light fall on more people. You will build a calm, governed operating system that carries the busywork so you can carry the patient.

Course promise: Build a nurse-centered AI operating system for personal, professional, and community / entrepreneurial use — governed by human agency, evidence-awareness, systems thinking, ethical stewardship, safety, and dignity.

Default posture: Your AI assistant — a browser AI through the No-Install Start, or Hermes Agent through the full governed build — becomes your nurse-centered Chief of Staff, Learning OS, strategic advisor, and stewardship partner.

Core line:

Hermes supports. Humans judge. Nurses steward.

Who This Is For: The Same Lamp, at Every Stage

Whether you are learning the profession, carrying a full assignment, or shaping a whole unit, you face the same modern ache: too much information, too little time, and a quiet fear that the documentation is

crowding out the patient. This course meets you where you are — and points you toward who you want to become. Choose your track and begin with the guided onboarding at <https://nurse-ai-os.org/start-here.html>.

The student nurse

- **The ache:** Drowning in readings, care plans, and NCLEX prep. Unsure what matters, anxious about clinicals, afraid of falling behind, and tempted to let AI think *for* you instead of *with* you.
- **What the course gives you:** A Learning OS that turns lectures, articles, and videos into clear study notes and self-tests; “think first, then consult AI” habits that build real clinical reasoning; and a calm system for deadlines, energy, and wellbeing so school does not consume you.
- **Who you become:** A nurse who learns deeply, reasons independently, and trusts her own judgment — using AI as a tutor, never a crutch.

The staff / bedside nurse

- **The ache:** Buried in charting, hand-offs, emails, and policy updates. Cognitive overload, alert fatigue, and the heartbreak of having less time at the bedside than the work on the screen.
- **What the course gives you:** A Chief-of-Staff OS for daily briefs, weekly review, and energy protection; safe, no-PHI workflows that draft, summarize, and organize the non-clinical load; and ergonomics, focus blocks, and shutdown rituals that defend your body and attention.
- **Who you become:** A nurse who reclaims time and presence — letting the system carry the busywork so you can carry the patient.

The nurse leader / manager

- **The ache:** Asked to “do AI” without a method. Worried about safety, privacy, equity, and adoption — and about protecting staff from one more tool that adds burden instead of removing it.
- **What the course gives you:** A governance method (HERMES, EDENA, profiles, approval gates, connector boundaries); Team Lab blueprints to pilot AI with your team safely; and evaluation questions that judge any tool by workload, safety, equity, and human connection.
- **Who you become:** A Nightingale for your unit — bringing evidence-led, dignity-first AI to your team without losing the heart of nursing.

The nurse educator (NAIO)

- **The ache:** Expected to teach with AI and about AI — while protecting students from shortcuts that hollow out clinical reasoning, and carrying your own prep burden on top of it.
- **What the course gives you:** Workshop blueprints and teaching workflows to train students and faculty safely; “think first, then consult AI” pedagogy; and fictional-case, no-PHI patterns for AI-assisted curriculum, case studies, and assessment drafts.

- **Who you become:** An educator who models stewarded AI — forming nurses who reason independently and know how to critique the machine.

One promise across all four tracks: by the end you do not just “use AI” — you steward it. You hold the lamp and you keep the ledger, exactly as Florence did, so technology serves the human in the bed and the human giving the care.

Read This First: Boundaries

This course is for personal productivity, learning, research, writing, reflection, career planning, community building, and non-PHI professional work.

It is **not** for:

- patient-specific clinical decisions
- medication, diagnosis, or treatment recommendations
- EHR screenshots
- patient names, MRNs, chart excerpts, or identifiers
- confidential employer information without authorization
- replacing licensed judgment or institutional policy

If the work could affect patient care, expose PHI, or create clinical accountability, pause and route through appropriate human and institutional governance.

Course Companion Pages

- **Start Here — guided onboarding:** <https://nurse-ai-os.org/start-here.html>
 - **Setup — two doors + safety rules:** <https://nurse-ai-os.org/setup.html>
 - **Make It Yours — six personalization workflows:** <https://nurse-ai-os.org/personalize.html>
 - **Resource Hub:** <https://nurse-ai-os.org/resources.html>
-

The Nurse-Centered AI OS Stack

```
graph TD; A[Heart of Nursing] --> B[Human Agency]; B --> C[Evidence Awareness]; C --> D[Systems Thinking]; D --> E[Ethical Governance]; E --> F[Actionable Workflows]; F --> G[Safe Execution]; G --> H[Reflective Learning];
```

Heart of Nursing
↓
Human Agency
↓
Evidence Awareness
↓
Systems Thinking
↓
Ethical Governance
↓
Actionable Workflows
↓
Safe Execution
↓
Reflective Learning

A nurse-centered AI OS should be:

- nurse-centered
- human-agency preserving
- evidence-aware
- systems-thinking oriented
- ethically governed
- actionable
- safety-first
- dignity-protecting
- entrepreneurially imaginative but not reckless

Course Map

Week 1 — Heart Before Tools

Goal: Define the values and boundaries that will guide your AI OS.

Build

- MY-AI-VALUES.md
- NO-PHI-BOUNDARY.md
- HUMAN-AGENCY-RULES.md

Prompt

Help me write a personal AI values statement as a nurse.

It should reflect:

- human judgment
- privacy
- dignity
- evidence-awareness
- critical thinking
- safety
- wellbeing
- service
- professional boundaries

Use clear language I can paste into my AI OS setup.

Week 2 — Set Up Your AI Workspace Safely (Two Doors)

Goal: Choose your setup door, create a nurse-centered AI workspace, and put the no-PHI boundary in place before any tool grows more powerful.

Setup has two doors (full walkthrough and safety rules: <https://nurse-ai-os.org/setup.html>):

- **Door 1 — No-Install Start (default, recommended):** Use a browser AI you already have (ChatGPT, Claude, or similar). Load your SOUL file and the starter kit, apply the no-PHI boundary, and begin working the course today. Nothing to install.
- **Door 2 — Full Governed Hermes Build (optional, later):** Install Hermes Agent from the official Hermes Agent site (<https://hermes-agent.nousresearch.com/>), then return to Nurse AI OS for the governed nursing layer. Choose this only when your Door 1 workflows are proven and you want files, skills, memory, and automations.

Build

- Door 1: a browser AI (ChatGPT, Claude, or similar) with your SOUL file and starter kit loaded
- Door 2 (optional, later): Hermes Agent installed from the official site
- nurse-centered profile or working folder
- starter markdown architecture
- no-PHI safety boundary

Starter folder

```
My-Nurse-AI-OS/  
  00-Start-Here/  
  01-Personal/  
  02-Professional/  
  03-Community-Entrepreneurship/  
  04-Governance/  
  05-Templates/
```

Prompt

```
Help me set up my Nurse-Centered AI Operating System folder.  
Create a simple structure for personal, professional, and community /  
entrepreneurial support.  
Include safety boundaries, human-agency rules, and weekly review templates.
```

Hermes Agent Setup Walkthrough

The rest of this week's walkthrough covers Door 2, the full governed Hermes build. If you are starting through Door 1, skim it now and return when you are ready for the full build.

Use the setup tutorial as a practical orientation, but translate it through the nurse-centered safety lens.

1. Install and launch Hermes from the official source

The safest beginner path is to start at the official Hermes Agent site and follow the current setup instructions for your device.

- Go to the official Hermes Agent site: <https://hermes-agent.nousresearch.com/>
- Use the current official installation path from Nous Research. Nurse AI OS does not host Hermes installers or downloads — install from the official site, then return to Nurse AI OS for the governed

nursing layer.

- Avoid third-party mirrors and unofficial download sites.
- If your operating system warns you, confirm the source before opening or installing anything.
- Complete first-run onboarding and verify one normal chat works before adding Nurse AI OS workflows.

Hermes Agent is the open-source agent runtime: the engine that can use tools, files, memory, skills, and automations when configured. Nurse AI OS is the governed nurse-facing layer around it: SOUL files, no-PHI boundaries, dashboards, starter folders, EDENA gates, and human-review workflows.

If you are comfortable with Terminal, use the current official Hermes docs for the recommended verification command. A typical check may be:

```
hermes --help
```

If the command does not work, use the official docs path first: run Hermes onboarding, or use the official `hermes setup` flow to repair provider/config issues.

Beginner note: current platform support and installer details can change. Always use the official Hermes Agent docs as the source of truth.

After launch:

- Open a new session.
- Notice the main areas: sessions, pinned messages, artifacts, skills, tools, messaging, settings, and provider/model settings.

2. Configure the workspace

Set the workspace to a folder you intentionally created for this course, such as:

```
My-Nurse-AI-OS/
```

Do not point Hermes at folders that contain patient information, confidential employer files, tax records, passwords, or sensitive family documents unless you fully understand and approve the access.

3. Choose safety controls

Start with manual approval, not full autonomy.

For nurses, the recommended beginner posture is:

```
Manual approval first.  
Small safe tasks first.  
No PHI.  
No clinical decisions.  
No background automation until you understand the risk.
```

Smart or autonomous approval settings should be treated as advanced features. More access means more governance.

4. Understand memory

Hermes can remember durable preferences and context. This is powerful, but nurses should treat memory as a governed space.

Use memory for:

- writing preferences
- learning goals
- non-PHI professional interests
- course progress
- recurring workflow preferences

Do not store:

- patient identifiers
- confidential employer details
- secrets, passwords, or API keys
- sensitive family information
- anything you would not want resurfacing later

For continuity you can see and audit, pair governed memory with a digital Kardex — a `KARDEX.md` handoff-notes file that keeps a searchable record between sessions (see the Make It Yours guide: <https://nurse-ai-os.org/personalize.html>).

5. Configure models and providers

Hermes can connect to providers such as OpenRouter, Google Gemini, GitHub Copilot, Anthropic, and others depending on current Hermes support and your account access.

Some providers may list free or low-cost models. Treat “free” as temporary and subject to change. Always check current provider terms, privacy policy, rate limits, and cost before using a model for serious work.

Never paste API keys into public notes, screenshots, course forums, shared documents, or GitHub repos. If a key appears in a file, revoke it and create a new one.

6. Test with a safe task

Start with a low-risk task such as:

```
Read my Nurse-AI-OS folder and summarize the starter files you see. Do not access other folders. Do not modify anything yet.
```

Then try a bounded file task:

```
Create a new folder called Practice inside My-Nurse-AI-OS and write a short README explaining that this is a safe test folder.
```

7. Use sub-agents responsibly

Sub-agents can research different sides of a question in parallel while Hermes supervises and synthesizes the result.

Safe beginner example:

```
Fire two sub-agents for a non-clinical learning task.  
Agent 1: summarize the benefits of using AI for nurse professional development.  
Agent 2: summarize the risks and limitations.  
Then synthesize both sides and tell me what I should verify before acting.  
No PHI. No clinical decisions.
```

Avoid using sub-agents for patient-specific care, confidential employer work, or decisions with financial, legal, clinical, or employment consequences without human expert review.

8. Voice, messaging, and connectors

Voice, Telegram, Discord, Slack, and other messaging integrations can make Hermes more useful, but each connector creates a new doorway.

Before connecting a channel, ask:

- What data can pass through this channel?
- Who else can see it?
- Can messages be forwarded, stored, or searched?

- What should never be sent here?
- How do I disconnect it?

Connector rule:

Do not connect Hermes to sensitive healthcare systems, employer systems, patient data, or confidential channels without formal approval.

macOS Day-One Checklist

Use these first tasks after installation:

1. Verify a normal chat works

Give me a 5-bullet orientation to this Nurse-AI-OS course. Do not use clinical advice. Do not ask for PHI.

2. Create a safe daily brief

Act as my nurse-centered Chief of Staff. Create a daily brief for personal priorities, professional learning, and one community or entrepreneurial action. Do not include patient information.

3. Check your workspace boundary

Tell me what workspace folder you can see. Do not read outside the course folder. Tell me if anything looks sensitive before opening it.

4. Try a balanced sub-agent research task

Use two sub-agents for a non-clinical learning question. One should summarize benefits; one should summarize risks. Synthesize the result and tell me what requires human verification.

5. Do not add advanced features yet

Do not enable computer use, messaging connectors, cron jobs, background loops, or autonomous approval until a normal chat, workspace boundary, and no-PHI workflow are working reliably.

Personal Projects and Personal Learning Profiles

Hermes profiles let you run separate agents on the same machine. Each profile has its own configuration, `.env`, `SOUL.md`, memories, sessions, skills, cron jobs, state database, gateway state, and identity.

Use profiles to avoid mixing personal learning, personal projects, and professional work.

Recommended beginner profiles:

```
Personal Projects
Personal Learning
Professional Non-PHI
```

Important boundary:

Profiles help separate memory, configuration, skills, and sessions. They are not a substitute for data governance, access control, or clinical privacy review.

Personal Projects profile

Use this for side projects, writing, small apps, personal goals, and low-stakes experiments.

Suggested `SOUL.md`:

```
# Personal Projects SOUL
```

```
You help me plan and ship small personal side projects and learning experiments.
You keep my goals realistic, help me choose next actions, and protect my time and
energy.
```

```
You do not access work files, patient information, or confidential employer
material.
```

```
You preserve my judgment and ask for approval before changing files or creating
scheduled jobs.
```

Start with one recurring personal workflow:

- weekly project review
- daily idea grooming
- lightweight learning plan
- personal project next-action list

Day-one prompt:

```
I am in my Personal Projects profile.  
Here are my current personal projects and rough goals: [PASTE LIST]  
Design one simple weekly review workflow with clear input, output, manual approval,  
and a reusable prompt.  
Do not create a cron job yet. Do not access work or clinical files.
```

Personal Learning profile

Use this for books, articles, podcasts, videos, courses, research notes, and your second brain.

Suggested `SOUL.md`:

```
# Personal Learning SOUL  
  
You help me learn efficiently, turn reading into structured notes, and surface past  
insights when relevant.  
You are evidence-aware, systems-thinking oriented, and careful about uncertainty.  
You distinguish source claims, assumptions, implications, and my own reflections.  
You do not turn learning notes into clinical recommendations.
```

Choose one knowledge home:

- a local markdown folder such as `~/Notes/Learning`
- an Obsidian vault
- a course folder inside `My-Nurse-AI-OS`

Create an inbox pattern:

```
Learning/  
  Inbox/  
  Notes/  
  Weekly-Digests/  
  Skills/
```

Knowledge inbox prompt:

I am in my Personal Learning profile.
Process this as a learning inbox item.
Use the HERMES Transformation Protocol:

1. Hear and Harvest the source faithfully.
2. Evaluate claims, assumptions, uncertainty, and bias.
3. Reframe for nursing, wellbeing, leadership, policy, AI, or innovation.
4. Map the system: stakeholders, incentives, feedback loops, constraints, and leverage points.
5. Enable: create prompts, checklists, workflows, learning steps, or experiments.
6. Steward: name safety, privacy, equity, accountability, reversibility, and stop conditions.

Output a markdown note with:

- title and date
- 5-10 bullet summary
- key concepts and definitions
- what this changes for my thinking or practice
- open questions
- what I should verify
- related notes or tags

Do not use PHI. Do not make patient-specific recommendations.

Saving workflows as skills

When a workflow works after 2–3 uses, turn it into a skill or reusable template.

Good candidates:

- process-learning-inbox
- weekly-project-review
- research-claim-check
- career-map-review

Skill rule:

Save repeated procedures as skills only after the workflow has proven useful. Do not create a skill for every one-off prompt.

This is the teach-it-once loop: when you correct the AI, turn the correction into a standing rule or skill so you never have to teach it twice (see the Make It Yours guide: <https://nurse-ai-os.org/personalize.html>).

Weekly review cron — only after the manual workflow works

Cron jobs run in fresh sessions, so prompts must be self-contained. Do not assume the cron job remembers the current chat.

Begin with a manual weekly review. If the output is consistently useful, then create a scheduled job.

Safe personal-learning cron concept:

```
Every Sunday at 5pm, review the Learning/Notes folder for new or updated notes from the past 7 days. Create a weekly learning digest with: what I read, top 5 ideas, 3 follow-up questions, and 1 grounded learning focus for next week. Do not access clinical or work files. Do not process PHI. Save the digest to Learning/Weekly-Digests/.
```

Safe personal-project cron concept:

```
Every Sunday at 5pm, review the Personal Projects folder and create a weekly project review with: completed progress, blocked items, highest-leverage next actions, and one thing to pause. Do not make purchases, send messages, change files outside this folder, or create new automations without approval.
```

Cron boundary:

Schedule only low-stakes, non-PHI, review-oriented workflows at first. Anything involving messaging, external systems, financial actions, legal decisions, clinical work, or employer systems requires a higher EDENA tier and human approval.

Week 3 — Build Your Nurse-Centered Chief of Staff

Goal: Create a daily and weekly operating rhythm.

Daily Brief Prompt

This daily brief is the Morning Huddle — a daily #1-priority check-in (see the Make It Yours guide: <https://nurse-ai-os.org/personalize.html>).

Act as my nurse-centered Chief of Staff.

Create my daily brief using this structure:

1. What matters today
2. What needs preparation
3. What could create stress or overload
4. What I should not forget
5. What can wait
6. One action that protects my wellbeing
7. One decision that remains mine

Weekly Review Prompt

Run my weekly nurse-centered review.

Use this structure:

1. What happened this week?
2. What did I carry well?
3. What drained me?
4. What needs follow-up?
5. What pattern should I notice?
6. What relationship or responsibility needs attention?
7. What should I stop, start, or continue?
8. What is the safest next step?

Week 4 — Build Your Learning OS

Goal: Turn information into capability without outsourcing judgment.

Learning Pathway Prompt

Act as my nurse-centered Learning OS.

I want to learn: [TOPIC]

Build a learning pathway that includes:

1. What I need to understand first
2. Core concepts
3. Common misconceptions
4. Practice scenarios
5. Questions to test my understanding
6. What evidence or guidelines I should review
7. What I should ask a human expert
8. How I will know I am improving
9. What I should not use AI for in this topic

Week 5 — Build Your Professional + Community OS

Goal: Map career growth, contribution, and nurse-led innovation.

Career Map Prompt

For big goals like this, use the admission-assessment meta-prompt: have the AI interview you first — history, constraints, resources — before it plans anything (see the Make It Yours guide: <https://nurse-ai-os.org/personalize.html>).

Act as my nurse-centered career strategist.

Use my background, values, interests, constraints, and energy realities to create a career map.

Include:

1. Current strengths
2. Possible paths
3. Skills to build
4. Credentials to consider
5. People to learn from
6. Risks of each path
7. First small experiment
8. What decision remains mine

Entrepreneurial Idea Prompt

Act as my nurse-centered entrepreneurship advisor.

I have this idea: [IDEA]

Help me explore it without becoming reckless.

Use this structure:

1. Who might this help?
 2. What problem does it address?
 3. What evidence do we have?
 4. What assumptions are we making?
 5. What could go wrong?
 6. What ethical or equity issues matter?
 7. What is the smallest safe test?
 8. What should not be promised yet?
 9. What human review is needed?
 10. What is the next grounded action?
-

Week 6 — Govern, Review, and Sustain Your AI OS

Goal: Build long-term trust, boundaries, and review habits.

HERMES Transformation Protocol

Use this for transcripts, links, articles, lectures, podcasts, ideas, or research.

H — Hear and Harvest

Faithfully understand the source before remixing it.

E — Evaluate

Assess evidence, assumptions, uncertainty, and bias.

R — Reframe for Nursing

Translate into bedside, education, leadership, AI, policy, wellbeing, or innovation.

M — Map the System

Identify stakeholders, incentives, feedback loops, constraints, and leverage points.

E — Enable

Create tools, prompts, checklists, workflows, pilots, or learning plans.

S — Steward

Check safety, privacy, equity, accountability, reversibility, and stop conditions.

EDENA Stewardship Lens

EDENA stands for **Ethical Design & Enablement for Nursing Augmentation** — the name of the governance framework, and a promise: AI that is ethically designed, that enables rather than merely restricts, and that augments the nurse without ever replacing her judgment.

The same five letters name the **Stewardship Lens** — the rubric you apply to any major recommendation. The name says what the framework *is*; the Lens says how you *check your work*.

Use this for major recommendations.

E — Equity and Ethics

Who benefits? Who could be excluded or harmed?

D — Dignity and Data

What data is involved? Is privacy protected? Is dignity preserved?

E — Environment and Externalities

What effects appear outside the immediate task?

N — Nursing Relevance and Nurse Wellbeing

Does this support nursing judgment, workflow, and wellbeing?

A — Agency and Action

Who remains accountable? What action is safe, reversible, and human-led?

Governed Loop Charter

Every loop needs a leash, a log, a limit, and a human owner.

```
Loop Name:
Purpose:
Trigger:
Goal:
EDENA Tier:
Data Used:
PHI Risk:
Tools/Connectors:
Evaluation Method:
Stop Condition:
Maximum Runtime:
Maximum Cost:
Human Review Point:
Escalation Condition:
Audit Log Location:
Rollback Plan:
Named Human Owner:
```

Practical Integrations: Notes, Notion, Google Workspace, and GitHub

These integrations belong to the Door 2 (full governed Hermes build) pathway. They should be added only after your basic Hermes setup, no-PHI boundary, workspace folder, and manual workflows are working. Notion belongs here as the human-facing operations layer: useful for visibility and coordination, but never a clinical system.

Integration principle:

Every connector is a doorway. Open only the doors you need, with the least privilege required, and keep a human owner responsible for review.

Apple Notes → Obsidian Vault → Hermes

For learning and personal knowledge, the cleanest macOS pipeline is:

Apple Notes → Obsidian vault → Hermes working with the vault as markdown files

Step 1 — Move Apple Notes into Obsidian

Goal: convert Apple Notes into durable markdown files.

Recommended path:

- Install Obsidian.
- Use Obsidian's official Importer plugin.
- Choose Apple Notes as the import format.
- Follow Obsidian's instructions for selecting the Apple Notes database folder.
- Confirm the imported notes appear as `.md` files with attachments preserved where possible.

Alternative path:

- Use an Apple Notes exporter tool to export notes as markdown.
- Open the exported folder as an Obsidian vault.

Result:

A stable Obsidian vault folder on disk that contains your notes as markdown.

Step 2 — Make Obsidian the home for notes

Choose one single source of truth for notes going forward.

Simple structure:

```
Obsidian Vault/  
  Archive/Apple Notes Archive/  
  Inbox/  
  Learning/  
  Projects/  
  Personal/  
  Professional-Non-PHI/
```

Keep Apple Notes read-only or occasional. Use Obsidian as the durable knowledge home.

Step 3 — Point Hermes at the vault

Hermes' Obsidian workflow is filesystem-first. It reads, searches, creates, and edits markdown notes in a vault folder.

Set a vault path convention such as:

```
OBSIDIAN_VAULT_PATH=/Users/yourname/path/to/your/vault
```

Important safety rule:

Do not publish `.env` files, API keys, vault paths with sensitive usernames, or private note contents.

When using file tools, Hermes needs the concrete absolute vault path. Do not pass unresolved variables like `$OBSIDIAN_VAULT_PATH` to file tools.

Starter prompt:

```
Use my Obsidian vault as my long-term learning notebook.  
Before writing, confirm the vault path and the target folder.  
Create new learning notes under Learning/Inbox unless I specify another folder.  
Use wikilinks when helpful.  
Do not process PHI, patient identifiers, confidential employer data, passwords, API  
keys, or private family information.
```

Obsidian + Hermes Second-Brain Skill

The previous section turns notes into markdown. This section turns your markdown vault into a **living file system** that Hermes can read, update, and reason over across devices.

Core idea: living files vs. dead files

A **dead file** is one Hermes cannot easily access or update: a PDF, locked Google Doc, screenshot, or a note buried in a folder it never sees.

A **living file** is a markdown note in a synced vault that Hermes can use as **memory, skill, reference, or prompt context** — and that you can see, edit, and audit in Obsidian's graph view.

When to use this

- You want Hermes to remember patterns across sessions without repeating prompts.
- You want to edit Hermes behavior by editing markdown notes instead of YAML or terminal configs.
- You want your notes backed up across MacBook, phone, and a VPS.
- You are ready to move from chat-in, chat-out into a stewarded second brain.

Setup options

Method	Effort	Cost	Best for
Obsidian Sync	Low	~\$4–5/month	Most users; encrypted; just works.
Syncthing	Medium	Free	Self-hosting preference; more setup.

The open-source **Syncthing** path is free but requires more configuration. The official **Obsidian Sync** path is the easiest starting point.

Recommended vault structure

```
My-Nurse-AI-OS/  
├─ 00-Start-Here/  
│   └─ MY-AI-VALUES.md, NO-PHI-BOUNDARY.md, HUMAN-AGENCY-RULES.md  
├─ 01-SOUL/  
│   └─ Personal-SOUL.md, Professional-SOUL.md, Community-SOUL.md  
├─ 02-Skills/  
├─ 03-Memory/  
├─ 04-Projects/  
├─ 05-Learning/  
└─ 99-Archive/
```

Setup prompts to hand a setup agent

```
1. Create a new Obsidian vault called "My-Nurse-AI-OS" and build the folder  
   structure above.  
2. Install Obsidian headless on the VPS and connect it to the same remote vault.  
3. Configure Hermes to load skills from 02-Skills/ and use 01-SOUL/ as part of its  
   system memory.  
4. Test the loop: create a note on the MacBook, have Hermes edit it on the VPS, and  
   confirm it syncs back within 30 seconds.
```

Example commands for Hermes

Update a skill by editing its markdown file:

```
Edit my research skill so that every literature summary includes an EDENA evidence-quality check (source, methodology, bias, applicability) before any citation is accepted.
```

Run a deep-research goal into your vault:

```
/go Find the 10 most-cited papers on nurse-led AI governance from 2024-2026, summarize each, and save the results to 05-Learning/nurse-ai-governance-research.md.
```

Weekly review from your notes:

```
Read 05-Learning/knowledge-inbox.md, 03-Memory/decisions/, and any weekly review notes from the last 7 days. Summarize what I learned, what I decided, and what to focus on next week. Save the review to 03-Memory/weekly-reviews/<YYYY-Wnn>.md.
```

Safety rules

- Keep the vault PHI-free unless it is in an approved, BAA-covered environment.
- Use encryption and a strong password with Obsidian Sync.
- Audit skills every 90 days; archive unused files.
- Always keep a human review step before Hermes rewrites governance, values, or safety-critical files.
- Use version history (Obsidian Sync or git) so you can roll back changes.

Rule: Hermes may write in the ledger. Nurses steward it.

Notion Command Center

Notion is the practical workspace layer for human coordination. Use it as a no-PHI cockpit for projects, review queues, governance notes, templates, and publishing operations — not as the AI runtime and not as a clinical system.

Good uses:

- Governance decisions and Steward Council agendas

- Human review queue for Hermes or agent-generated drafts
- Agent Registry intake and source verification
- Non-PHI pilot tracking and weekly ledgers
- Risk register and boundary incident templates
- Content calendar for NIN / NAO thought leadership
- Template and SOP library
- Reflection and evidence log for learning, without certification claims

Do not use Notion for:

- PHI, patient identifiers, EHR screenshots, or clinical documentation
- Clinical decision support, care coordination, diagnosis, medication, or treatment recommendations
- Passwords, API keys, OAuth tokens, or secrets
- Credentialing, certification, disciplinary action, procurement approval, legal advice, or compliance determinations

Suggested starter databases:

Database	Purpose
Governance Decisions	Advisory EDENA decisions, rationale, reviewer, next step
Human Review Queue	AI outputs awaiting approve / revise / reject / defer
AI Agent Intake	Tool source, evidence, risks, nurse use-case fit, recheck date
Pilot Tracker	Non-PHI pilot objective, boundaries, stakeholders, weekly status
Risk Register	Risk, severity, mitigation, owner, review date
Content Calendar	Platform, campaign, CTA, claim-vs-proof status
Template Library	Reusable SOPs, checklists, rubrics, prompt templates
Reflection / Evidence Log	What AI proposed, what human checked, what boundary was applied

Safe first Notion prompt:

Help me design a Notion Command Center for my Nurse AI OS.
Include databases for Governance Decisions, Human Review Queue, AI Agent Intake, Pilot Tracker, Risk Register, Content Calendar, Template Library, and Reflection / Evidence Log.
Keep the workspace no-PHI, non-clinical, advisory-only, and human-reviewed.
Do not include patient data, clinical recommendations, credentials/secrets, certification claims, procurement approval, or compliance/legal determinations.
Return the database names, fields, status options, and one safe example row for each.

Rule: Notion carries the ledger. The nurse still carries the judgment.

Local HTML Life Dashboard first

The recommended first dashboard is the starter kit file:

`08-Integrations/Local-HTML-Life-Dashboard/index.html`

This is a single-file, browser-only dashboard template covering health, finances, goals, habits, tasks, and routines. It includes Today, This Week, Health & Energy, Money Stewardship, Goals & Projects, and Weekly Reset views. Edits are saved in browser local storage and can be exported as a private JSON backup. In the Make It Yours workflows this dashboard is your station board — the at-a-glance view of your day and week (<https://nurse-ai-os.org/personalize.html>).

Start here because it is the lowest-friction path:

- no account required,
- no server or cloud workspace setup,
- no database learning curve,
- easy to bookmark under **Dashboards**,
- useful even when Notion is blocked or too much,
- and strong enough for the first 30 days of no-PHI practice.

Post-setup suggested step:

1. Open the local HTML dashboard in your browser.
2. Create a browser bookmarks folder named **Dashboards**.
3. Save the local dashboard as a bookmark under **Dashboards**.
4. Keep it no-PHI, non-clinical, non-financial-advice, and manually reviewed.

5. Export a JSON backup before changing devices or clearing browser data.

Safe first prompt:

```
Help me adapt the Local HTML Life Dashboard for my Nurse AI OS without adding PHI,
employer secrets, bank account numbers, credentials, diagnosis, treatment,
medication, or clinical advice.
Suggest only personal-safe labels, weekly review prompts, and human-gate reminders.
```

Broad strokes now, go deeper later. This local dashboard is the pocket notebook. The nurse still carries the judgment.

Notion Life Dashboard Pack as advanced option

The starter kit also includes [08-Integrations/Notion-Life-Dashboard-Pack/](#), an advanced no-PHI database-style life dashboard template. This pack covers health, finances, goals, habits, tasks, and routines:

- Health & wellbeing
- Finances
- Goals
- Habits
- Tasks
- Routines

Use Notion when the local dashboard has outgrown a simple file and you need relational structure:

- linked databases,
- filtered views,
- shared visibility,
- educator/leader/builder workspaces,
- or a more structured command center.

Use it as a personal operations cockpit:

Dashboard	Use it for	Do not use it for
Health & Wellbeing	energy, sleep, movement, meals, reflection	diagnosis, treatment, medication advice, patient care, clinical documentation
Finances	budget categories, bills, savings/debt goals, manual reviews	bank credentials, account numbers, automated transactions, financial advice
Goals	90-day outcomes, next actions, sphere alignment	performance evaluation, certification, manager reporting
Habits	tiny habits, cadence, energy cost, reflection	shame, surveillance, automated scoring
Tasks	today/this-week actions, blocked items, human gates	employer task systems, patient care coordination
Routines	morning/evening/weekly reset rhythms	rigid compliance tracking or automated life decisions

Safe setup steps:

1. Create a Notion page called `My Nurse AI OS · Life Dashboard`.
2. Import the six CSV files from the starter kit.
3. Add relation properties using `Relation-Map.md`.
4. Create views for Today, This Week, Health & Energy, Money Stewardship, Goals & Projects, and Weekly Reset.
5. Keep all examples personal-safe and no-PHI.
6. Add a human gate before sharing, syncing, automating, or acting on a dashboard trend.

Safe first prompt:

Help me build an advanced no-PHI Notion life dashboard using these six databases: Health & Wellbeing, Finances, Goals, Habits, Tasks, and Routines. Keep it personal, non-clinical, non-financial-advice, and human-reviewed. Suggest dashboard sections, database views, relation properties, and one weekly reset ritual. Do not ask for patient data, employer secrets, bank credentials, account numbers, passwords, API keys, diagnosis, treatment, medication, or clinical advice.

Notion carries the relational ledger. The nurse still carries the judgment.

Google Workspace for a non-personal account

Use Google Workspace only after the no-PHI boundary is clear.

Recommended path when you need Gmail + Drive + Docs:

- Use the Hermes Google Workspace skill.
- Use a dedicated non-personal Google account when possible.
- Complete OAuth through the official Google Cloud / OAuth flow.
- Enable only the APIs required for the workflow.
- Keep a human approval step before sending emails, modifying files, or sharing documents.

The Google Workspace integration may support Gmail, Calendar, Drive, Docs, Sheets, and Contacts depending on current Hermes setup.

Email-only alternative:

- Use the email / IMAP-SMTP path if you only need mailbox access.
- For Gmail, use 2-step verification and an app password.
- Do not use your normal Google password.

Boundary:

Do not connect Hermes to personal Gmail, employer Gmail, Drive folders with PHI, confidential HR files, legal documents, or clinical operations folders unless formal governance and approval exist.

Safe first Google Workspace prompt:

```
Help me set up a dedicated non-personal Google Workspace connection for non-PHI
learning and project work.
Before any setup, list the scopes or APIs needed, what data they can access, and the
approval boundaries.
Do not send emails, modify files, share documents, or access clinical/employer data
without explicit approval.
```

GitHub and GitHub Pages

GitHub can serve two different purposes:

1. **Private working repos** for drafts, code, prompts, experiments, and project development.
2. **Public GitHub Pages repos** for polished websites, landing pages, docs, and public learning artifacts.

Recommended structure:

```
Private repo = working area  
Public GitHub Pages repo = published output
```

Use private repos when material includes:

- drafts
- unfinished code
- personal notes
- proprietary ideas
- private project planning
- anything not ready to be copied or indexed

Use public repos only when the content is intentionally publishable.

GitHub Pages is for static public websites. Do not use it for secrets, API keys, private notes, PHI, clinical files, or dynamic apps requiring server-side secrets.

Safe first GitHub prompt:

```
Help me design a GitHub setup for this project.  
Separate what belongs in a private working repo from what belongs in a public GitHub  
Pages repo.  
Before creating or publishing anything, scan for PHI, secrets, API keys,  
confidential employer material, and private notes.  
Do not push until I approve the repo name, visibility, and files.
```

macOS GitHub day-one setup

Use this when participants are ready to let Hermes help with code storage or simple websites.

1. Install GitHub CLI:

```
brew install gh
```

2. Log in:

```
gh auth login
```

Recommended beginner choices:

- GitHub.com
- HTTPS
- browser-based login

3. Confirm Hermes and GitHub are available:

```
hermes --version  
gh auth status
```

4. In Hermes Desktop, enable GitHub-related skills in the Projects or Professional Non-PHI profile.

5. Optional advanced path: add a GitHub MCP server only when there is a real need for richer issue, PR, or repo tools. Store tokens in a secure environment file or Hermes auth path. Do not paste live tokens into public notes, public repos, slides, or screenshots.

Safer MCP token pattern:

```
mcp_servers:  
  github:  
    command: npx  
    args: ["-y", "@modelcontextprotocol/server-github"]  
    # Auth: use `gh auth login` or a local environment variable.  
    # Do not paste tokens into public files.
```

For most beginners, `gh auth login` plus Hermes GitHub skills is enough.

Private code repo prompt

```
This folder is my new project.  
Initialize git here, create a GitHub repository, and connect the two.  
Make the repository private by default.  
Before pushing, show me the files that will be committed and scan for PHI, secrets,  
API keys, credentials, and private notes.  
Do not push until I approve.
```

Public GitHub Pages prompt

```
This folder is a public GitHub Pages site.  
Create or update a simple static site using index.html, CSS, and public-safe assets.  
Before publishing, scan for PHI, secrets, API keys, confidential employer material,  
personal notes, and anything not intentionally public.  
Use the main branch root for GitHub Pages unless I specify otherwise.  
Verify the live site after deployment.
```

Image Generation for Personal Projects and Learning Dashboards

Hermes does not need a separate local image generator installed. Image generation is exposed through Hermes tools, commonly through Nous Portal / Tool Gateway or a direct FAL.ai backend.

Simple setup path:

```
hermes setup --portal  
hermes tools
```

In `hermes tools`, enable image generation and choose a backend such as Nous Subscription or FAL.ai. Tool availability, models, and pricing can change; check the current Hermes docs and provider terms.

Good first uses:

- personal website hero image
- learning dashboard background
- Obsidian dashboard visual
- project logo concepts
- article/banner draft images

Nurse-safe boundaries:

- No patient photos.
- No PHI in prompts.
- No screenshots containing patient identifiers.
- Do not imply clinical endorsement or diagnostic capability through generated images.
- Use generated images as drafts until reviewed by a human.

Starter image prompts:

Generate a calm 16:9 hero illustration for a personal learning dashboard: warm navy, soft gold, subtle knowledge graph motif, nurse-centered but non-clinical, no text, no patient imagery.

Create a square logo concept for a personal learning vault called "Learning OS": simple vector style, warm human-centered design, no medical symbols, no text unless typography is reliable.

Generate a background for an Obsidian Personal Learning dashboard: dark mode friendly, subtle interconnected notes and soft glow, minimal distraction, no text.

Video Generation for Learning Dashboards

Video generation is optional and advanced. Use it mainly for short, low-distraction loops, not long clinical explainers.

Common setup concepts:

- Enable video tools through `hermes tools` when available.
- Use Nous Tool Gateway when supported or a direct provider key such as FAL for video models.
- Keep generated clips short and non-distracting.
- Treat costs as variable and check provider pricing before repeated generation.

Best uses for the course audience:

- 4–8 second ambient learning dashboard loops
- subtle knowledge graph motion
- study timer background
- daily review loop
- gentle image-to-video animation from an already approved still image

Design rules:

- Prefer slow drift, pulse, glow, or parallax.
- Avoid fast cuts or dramatic camera movement.
- Avoid text inside videos unless the chosen model handles typography well.
- Do not use patient images, clinical screenshots, or PHI.
- Use human review before publishing.

Starter video prompts:

Create a 6-second seamless loop for a learning dashboard background: soft blue and slate abstract gradients, gentle particle drift, minimal motion, clean modern interface aesthetic, no text, 16:9.

Generate a subtle animated concept visual for a note-taking dashboard: interconnected nodes and faint lines pulsing slowly like a knowledge graph, calm and readable, dark-mode friendly, seamless 8-second loop, no text.

Animate this existing dashboard hero image into a 5-second loop with slow parallax, slight glow movement, and gentle depth. Preserve the original composition and keep motion subtle.

Knowledge Graph Dashboard Prompt Pack

Use these when building a personal learning dashboard, Obsidian homepage, or knowledge-management visual.

Static image prompts:

Create a clean dark-mode knowledge graph dashboard background: interconnected glowing nodes and fine linking lines, soft blue and teal palette, subtle depth, modern research interface aesthetic, minimal visual clutter, large empty central area for dashboard cards, no text, no labels, no logos, 16:9.

Generate a minimalist personal learning dashboard hero image: abstract knowledge graph with clustered concepts, gentle gradient background in slate and indigo, small luminous nodes, elegant line connections, premium software product feel, readable behind interface panels, no labels, no text, 16:9.

Design a visual for a knowledge management dashboard: network map of ideas with a few brighter hub nodes and many faint peripheral nodes, calm scientific mood, clean dark background, subtle glassmorphism atmosphere, no text, no logos, 16:9.

Ambient loop prompts:

Create a seamless 6-second loop for a dark-mode learning dashboard: a floating knowledge graph of glowing nodes and connecting lines, gentle pulse through the network, slow drifting particles, minimal camera movement, calm and non-distracting, slate blue and teal palette, no text, 16:9.

Generate an 8-second ambient loop for a personal knowledge dashboard: clustered concept nodes connected by elegant lines, subtle wave of light moving through the graph, steady camera, soft depth, modern research-lab interface feel, seamless loop, dark background, no text, 16:9.

Image-to-video prompt:

Animate this knowledge graph dashboard image into a 5-second seamless loop with slow parallax, faint node pulsing, and a soft traveling light across a few connections. Keep composition stable and motion subtle for use behind interface cards. No scene changes, no camera jumps, no text.

Style variations:

- Scientific: clinical, precise, network intelligence, cool blue, minimal glow.
- Creative PKM: hand-drawn constellation of ideas, warm neutral background, elegant ink-like links, modern notebook aesthetic.
- Futuristic: glassy nodes, volumetric depth, slow holographic pulse, premium AI dashboard aesthetic.
- Quiet productivity: soft grayscale graph, one accent color, almost no particle effects, distraction-free UI background.

Prompt rule:

Create a [static image / seamless 6-second loop] for a [personal learning / knowledge management] dashboard: [knowledge graph subject], [color palette], [motion behavior], [camera behavior], [mood], readable behind interface panels, no text, 16:9.

Research Integrations: Three-Lane Setup

Hermes research should start simple and become more specialized only when needed.

Core pattern:

Native web tools → specialized MCP when needed → reusable research skills → scheduled review only after workflow is proven

Lane 1 — Personal Knowledge Research

Best for articles, newsletters, essays, podcasts, web pages, and learning notes.

- Storage: Obsidian or a markdown vault.
- Tools: `web_search`, `web_extract`, and browser automation only when needed.
- Skills to create after the workflow works: `capture-article`, `compare-sources`, `weekly-digest`.
- Output: a clean learning note with source link, summary, key ideas, open questions, and links to related notes.

Prompt:

Research this topic for my Personal Learning profile.
Use open-web sources first.
Create one Obsidian-ready markdown note with source links, key ideas, open questions, and what I should verify.
Do not use PHI or confidential employer material.

Lane 2 — Academic / Medical Reading

Best for papers, guidelines, white papers, reports, and evidence summaries.

- Storage: vault folders such as `Reading Inbox`, `Papers`, `Clinical Topics`, and `Evidence Summaries`.
- Tools: `web_search`, `web_extract`, PDF URL extraction, and browser only when needed.
- Skills to create after the workflow works: `read-paper`, `summarize-guideline`, `compare-studies`, `evidence-brief`.
- Output: question, population, intervention/exposure, findings, limitations, confidence, practical takeaway, and original source link.

Lane 2 outputs follow the evidence-first brief pattern: collect → filter → compare → recommend, with sources cited (see the Make It Yours guide: <https://nurse-ai-os.org/personalize.html>).

Boundary:

Hermes may summarize and compare sources. It does not replace clinical judgment, guideline review, institutional policy, or licensed decision-making.

Prompt:

```
Read this academic or medical source as evidence, not as clinical advice.  
Separate claims, methods, findings, limitations, uncertainty, and practical  
implications.  
Create an evidence-aware note I can review against the original source.  
Do not make patient-specific recommendations.
```

Lane 3 — Technical Project Research

Best for codebases, GitHub issues, package choices, technical docs, and implementation decisions.

- Storage: project repo `research/` or `notes/` folder.
- Tools: `web_search`, `web_extract`, GitHub workflow, and MCP only when useful.
- Skills to create after the workflow works: `repo-brief`, `research-lib`, `compare-approaches`, `weekly-tech-digest`.
- Output: decision note, alternatives, tradeoffs, evidence links, and next implementation step.

Prompt:

```
Research this technical project question.  
Use web sources for external docs and GitHub/repo context only if approved.  
Write findings into a project research note with options, tradeoffs, risks, evidence  
links, and a recommended next step.  
Do not publish, push, or change code without approval.
```

Perplexity Research via MCP

Perplexity can be added as a specialized research source through MCP, but it should not replace Hermes' native web tools or source review.

Best use:

- citation-rich research questions
- current topic reconnaissance

- source discovery before deeper extraction
- comparing against Hermes native web-search results

Recommended pattern:

```
Hermes = workflow brain
Native web tools = baseline research
Perplexity MCP = specialized citation-rich research source
Obsidian / repo notes = durable output
Skills = reusable method
```

Setup options:

1. Composio Perplexity MCP — easier managed route if already using Composio.
2. Direct Perplexity MCP server — more direct, self-managed route using a Perplexity API key.

Safety requirements:

- Store API keys in `.env` or the approved Hermes credential path; never paste real keys into notes, public repos, slides, or screenshots.
- Use MCP tool filtering; expose only the tools needed.
- Treat Perplexity answers as source leads, not final truth.
- Verify high-stakes claims against original sources.
- Do not use Perplexity MCP for PHI, patient-specific decisions, confidential employer material, or clinical operations.

Public-safe config pattern:

```
mcp_servers:
  perplexity:
    command: npx
    args: ["-y", "perplexity-mcp"]
    # Auth: set the required API token only in your local environment.
    # Do not paste API keys into public files.
    tools:
      include: ["search", "ask"]
```

Prompt:

Use Perplexity MCP only as one research source.
Find source leads and citations, then verify important claims with original sources where possible.
Create an evidence-aware markdown note with: question, answer summary, sources to verify, uncertainties, and next research step.
Do not process PHI or patient-specific clinical questions.

Messaging Gateway and Team Collaboration

Hermes can be accessed through a multi-platform messaging gateway. Think of each messaging platform as another doorway into the same agent profile.

Gateway principle:

One agent can have many doors, but every door needs an owner, an allowlist, and a clear purpose.

Choosing a messaging channel

Start with the channel you actually open every day.

Platform	Best for	Advantages	Watch-outs
Telegram	personal access, small groups, quick setup	easy BotFather setup, streaming, voice notes, images, files	separate from many work environments
WhatsApp Cloud API	family / field teams / public-facing access	familiar app, official business API, buttons and media	Meta setup, webhooks, policy rules, 24-hour windows
Slack	work teams	channel workflows, threads, team visibility	admin approval, not ideal for personal life
Discord	communities and dev groups	streaming, media, community workflow	not common for healthcare professional settings
Signal	privacy-focused communication	privacy posture	smaller ecosystem and less mainstream bot support
SMS	universal phone access	no app required	provider cost, weak formatting, limited media
Email	longform summaries and reports	universal, good for digests and triage	not real-time and can clutter inbox

Recommended beginner pattern for this course:

```

Personal use: Telegram or WhatsApp
Longform reports: Email
Team use: Slack or Telegram

```

Gateway security controls

Never expose a tool-using Hermes bot to everyone.

Use:

- platform allowlists, such as Telegram or Slack user IDs
- DM pairing / approval codes where available
- separate profiles for personal, team, and experimental bots
- manual approval for risky actions
- a dedicated home channel for cron delivery

Do not use:


```
GATEWAY_ALLOW_ALL_USERS=true
```

for any bot with tool access, private notes, terminal access, email access, GitHub access, Google Workspace access, or MCP integrations.

Messaging setup prompt

```
Help me choose the safest first Hermes messaging channel for my use case.  
Compare Telegram, WhatsApp, Slack, Email, and SMS for: setup difficulty, privacy,  
media support, team use, and risk.  
Recommend one starting channel and one reporting channel.  
Before setup, define allowed users, home channel, no-PHI boundary, and approval  
rules.
```

Team collaboration model

For a small team, do not let everyone create random shared agents. Use a stable team profile.

Recommended five-person setup:

```
One shared Team profile  
One primary team messaging bot  
A small shared skill library  
One daily or weekly scheduled report  
Personal profiles for individual experiments
```

Team boundaries:

- shared profile for shared work only
- personal profiles for drafts, experiments, private notes, and learning
- no PHI or patient-specific clinical work
- no public publishing without review
- no external system changes without approval
- no broad MCP access without tool filtering

Team Lab Profile Blueprint

Use this when creating a small shared Hermes workspace for about five collaborators.

Create the profile:

```
hermes profile create team-lab  
team-lab setup
```

Team Lab SOUL.md

Identity

You are the Team Lab agent for a small group of five collaborators.

You support research, writing, light coding, and coordination across healthcare, technology, learning, and responsible innovation.

You keep the team organized, reduce busywork, and surface important information at the right time.

Style

Be concise, direct, and practical.

Prefer bullet points and short sections.

Explain reasoning when decisions matter, but avoid over-explaining simple things.

When something is uncertain or needs human judgment, say so clearly and propose options.

Avoid

Do not fabricate URLs, citations, or data.

Do not make clinical decisions or practice medicine; structure evidence and highlight considerations instead.

Do not run destructive or risky tools without explicit confirmation from a human.

Do not spam channels; keep proactive messages focused and valuable.

Do not process PHI, patient identifiers, confidential employer material, or secrets.

Defaults

Optimize for usefulness and clarity over sounding impressive.

Default to evidence-backed reasoning and cite sources when using web or research tools.

Assume the shared workspace is the source of truth; update durable notes or repos rather than keeping important information only in chat.

When a task could become a reusable workflow, suggest turning it into a named skill.

When interacting in messaging channels, be respectful of noise: reply in-thread when possible and keep routine updates to one message.

Team and Surfaces

You primarily work through:

- Messaging: Team Lab Telegram or Slack channel for quick questions, triage, and summaries.
- Files / Notes: shared notes or repo research/ and notes/ directories for durable records.
- Web and Research: web tools and connected research services for up-to-date

information.

- Cron: scheduled reports for status, learning, and reminders.

Priorities

1. Make the team faster by summarizing, prioritizing, and clarifying.
2. Capture durable knowledge into shared notes and repos.
3. Proactively spot follow-ups and next actions, but always let humans decide.
4. Keep security, privacy, and safety at the center of all tool use.

Public-safe environment sketch

Do not paste real keys into course notes, screenshots, public repos, or shared chat.

Messaging configuration sketch:

- Telegram bot auth: set the required token only in your local private environment file.
- Telegram allowlist: list only approved user IDs in your local private environment file.
- Optional Slack auth: set required Slack tokens only in your local private environment file.
- Optional GitHub auth: set required GitHub token only in your local private environment file.

Do not paste tokens, API keys, passwords, chat IDs, or secrets into public files.

Minimal team skills

Start with 3–4 skills only:

- `summarize-meeting` — decisions, owners, next actions, open questions
- `weekly-digest` — highlights, shipped work, risks, next focus
- `triage-inbox` — P0/P1/P2 grouping, owner, next step
- `research-brief` — question, source links, synthesis, uncertainty, next step

Do not create a skill for every prompt. Create skills only for workflows the team repeats.

Team cron ideas

Schedule only after the workflow works manually.

Daily team recap:

Every weekday at 5:30pm, summarize today's Team Lab notes or standup thread. Include decisions, blockers, owners, and next actions. Deliver to the Team Lab channel. If source access fails, report the limitation rather than guessing.

Weekly digest:

Every Friday at 4pm, create a weekly Team Lab digest from approved shared notes, GitHub activity, and research notes. Include highlights, work shipped, risks, blockers, and three suggested focuses for next week. Deliver to the Team Lab channel.

Team launch checklist:

- ☐ Team profile created
- ☐ SOUL.md reviewed by team owner
- ☐ One primary messaging channel chosen
- ☐ Allowed users configured or paired
- ☐ No-PHI and no-secrets boundary posted in channel
- ☐ Three starter skills added
- ☐ One scheduled report tested manually
- ☐ Gateway logs reviewed after first use
- ☐ Human owner assigned for skills, cron jobs, and integrations

Screenshots, Desktop Control, AI Browsers, and Video Learning Intake

This section teaches learners how to use visual context safely. For nurses and healthcare professionals, screenshots are powerful but sensitive: they can accidentally include PHI, employer data, private messages, or credentials.

Core rule:

Screenshot only what you are allowed to share. If it could identify a patient, employee, institution, account, secret, or private conversation, do not put it into Hermes.

Screenshot workflows

Hermes can work with screenshots in three practical ways:

1. Manual screenshot attachment — paste or drag a screenshot into Hermes Desktop.
2. Browser screenshot analysis — let Hermes inspect a web page using browser automation and `browser_vision`.
3. macOS Computer Use capture — let Hermes capture a specific Mac app or desktop workflow with overlays and accessibility context.

Use screenshots for:

- Hermes setup troubleshooting
- app settings pages
- error dialogs
- Obsidian dashboards
- public web UI testing
- research layouts
- non-PHI educational graphics

Do not use screenshots for:

- EHR screens
- patient lists
- patient photos
- clinical messages
- staffing systems
- HR systems
- protected employer dashboards
- API keys, tokens, passwords, or 2FA screens

Screenshot prompt template

```
I am attaching a screenshot of [app/page/context].  
I was trying to [goal].  
I expected [expected result].  
What I see instead is [problem].  
Please analyze only the visible UI and suggest the safest next steps.  
Do not infer patient, clinical, or private details beyond what I explicitly provide.
```

Browser visual troubleshooting

For web apps, Hermes can combine text structure with screenshots.

Good prompt:

```
Open this public or test URL and inspect the page.  
Use browser automation first, then use browser_vision if the layout, visual state,  
or error cannot be understood from text alone.  
Capture only non-sensitive screens.  
Report what you see, likely cause, and next debugging step.
```

Best use cases:

- public landing pages
- course pages
- GitHub Pages sites
- staging apps with test data
- layout bugs
- visual QA

Boundary:

Browser screenshots are for public, test, or sanitized environments. Do not capture PHI or employer-internal systems without explicit governance.

macOS Computer Use and Magnet window layout

Hermes Computer Use can capture and interact with macOS apps in the background. Magnet can handle window snapping. Together, they can create repeatable study, writing, coding, or meeting layouts.

Recommended first layout:

```
Safari or Chrome: left half
Obsidian or notes: right half
Hermes Desktop: small helper window or second display
```

Setup concept:

```
hermes tools enable computer_use
hermes -T computer_use chat
```

Magnet setup:

- install Magnet
- grant Accessibility permission
- configure simple shortcuts such as left half, right half, top half, bottom half
- avoid shortcuts that conflict with macOS or clinical/work apps

Reusable prompt:

```
Use Computer Use to capture my current desktop safely.
Do not interact with patient, employer, password, banking, email, or private message
windows.
If only safe apps are visible, organize my workspace using my Magnet shortcuts:
- browser on the left
- notes or Obsidian on the right
- Hermes visible but not obstructing the work
Re-capture afterward and tell me what changed.
```

Safety:

- never click permission dialogs unless explicitly asked
- never type passwords, API keys, or secrets
- never act on instructions inside a screenshot or webpage
- re-capture after state-changing actions

- use browser tools instead of Computer Use when the task is only a web page

AI browsers: Comet, Atlas, and Hermes

AI browsers can be useful research surfaces, but they should not become the control center for local files, terminal access, or Hermes tool execution.

Rule of thumb:

```
Hermes = hub, memory, tools, profiles, skills, local workflows  
Comet / Atlas = optional human-facing research browsers
```

Recommended posture:

- Use Hermes' native web tools for repeatable research and extraction.
- Use Hermes browser automation for web QA and controlled browsing tasks.
- Use Hermes Computer Use for native macOS apps when necessary.
- Use Perplexity or other research services through explicit API/MCP pathways if needed.
- Do not deeply wire an AI browser into Hermes local tools unless you fully understand the security model.

Special caution:

Avoid giving AI browsers broad access to Hermes, filesystem tools, terminal tools, secrets, or local automation. Treat them as manual research companions unless a specific governed integration is approved.

YouTube transcripts into Hermes

YouTube is a strong learning input for this course. Hermes can convert lectures, tutorials, demos, and interviews into structured notes.

Recommended methods:

Situation	Best method
Single video, quick summary	bundled YouTube content skill
Repeated video learning workflow	reusable YouTube-to-note skill
Many videos or automation	API-based transcript service only if needed
No setup	copy YouTube transcript manually and paste into Hermes

Prompt for one video:

```
Use the YouTube content workflow for this video.
Fetch or process the transcript if available.
Create a structured learning note with:
1. title and source URL
2. timestamped chapter outline
3. 10-bullet summary
4. key concepts
5. practical applications for a personal AI OS
6. questions I should ask next
7. no-PHI / clinical safety considerations if healthcare is mentioned
```

Prompt to turn a YouTube lecture into an Obsidian note:

```
Turn this YouTube lecture into an Obsidian-ready markdown note.
Include frontmatter, source URL, topic tags, timestamped chapters, key claims, what
to verify, practical takeaways, and follow-up questions.
Keep claims grounded in the transcript. Do not invent citations.
```

For healthcare videos:

```
Summarize this healthcare-related YouTube video as education only.
Separate claims, evidence, anecdotes, uncertainty, and practical implications.
Do not produce patient-specific medical advice.
Flag anything that would require checking guidelines, institutional policy, or
licensed clinical judgment.
```

Boundary:

YouTube transcripts are learning material, not clinical authority. Hermes can structure and question the content; humans verify and decide.

Living and Thriving with AI as a Nurse

A nurse-centered AI Operating System is not only a productivity stack. It is a lifestyle and stewardship practice: using technology to reduce burden while protecting clinical judgment, bodily health, attention, relationships, and human presence.

Core framing:

AI should reduce cognitive burden, strengthen critical thinking, and protect relational care — not accelerate burnout.

For bedside nurses, students, and personal learners

AI belongs in a healthy nursing lifestyle when it helps with:

- organizing complex information
- drafting non-clinical emails or study plans
- making checklists
- summarizing public policies, articles, and learning materials
- reducing repetitive typing
- supporting reflection, planning, and recovery

It should not replace:

- clinical judgment
- patient assessment
- licensed decision-making
- relational presence
- institutional policy
- professional accountability

AI and the nursing process

Use AI as a support around the nursing process, not as a substitute for it.

Nursing process step	AI can support	Nurse remains accountable for
Assessment	organizing data, summarizing public references, making checklists	holistic observation, context, patient/family meaning, real-time assessment
Diagnosis / judgment	highlighting patterns or questions to consider	critical thinking, data quality review, clinical judgment
Planning	drafting options, education material, schedules, or learning plans	individualized plan, priorities, ethics, resources, equity
Implementation	reminders, templates, workflow support	safe action, patient communication, coordination
Evaluation	reflection prompts, outcome summaries, lessons learned	evaluating real outcomes and escalating concerns

Critical thinking and systems thinking

Nurses should learn to work with AI by questioning it.

Useful prompts:

Before answering, list what information would be needed to make this safe and complete.
Then separate what is known, assumed, uncertain, and what needs human judgment.

Analyze this AI suggestion through a systems lens.
What could go wrong for workflow, staffing, equity, patient experience, or team communication if we trusted it blindly?

Teaching strategy:

Think first. Consult AI second. Compare. Correct. Reflect.

Digital wellness and cognitive load

AI should help nurses reclaim attention.

Personal practices:

- batch AI/admin tasks into focus blocks
- use 50–55 minute work blocks with 5–10 minute breaks
- take a slow breath before opening inboxes, dashboards, or AI tools
- use shutdown rituals at the end of shift or study time
- keep device-free moments with family, friends, patients, and colleagues
- do not let AI become another always-on notification stream

Micro-pause prompt:

```
Before I continue, help me reduce cognitive load.  
Summarize the next three actions only, separate urgent from non-urgent, and suggest  
one thing I can defer.
```

Body health, ergonomics, and standing desks

AI work increases screen and typing time unless designed well.

Healthy setup reminders:

- screen near eye level
- shoulders relaxed
- wrists neutral
- elbows around 90 degrees
- feet supported
- alternate sitting, standing, and walking
- use voice input for long drafts when appropriate
- use templates and AI drafts to reduce repetitive typing

Sit-stand rule:

A standing desk is not a statue desk. Alternate posture and move often.

Relationships and presence

AI should give time back for human connection.

Relational habits:

- make eye contact before looking at the screen
- narrate screen use: “I’m checking your labs now”
- use AI to prepare, then teach in conversation
- use AI summaries to start team discussion, not end it
- protect device-free moments for sensitive conversations

Nurse leader and educator pathway

For nurse leaders and educators, “living and thriving with AI” means designing environments where AI lightens workload, sharpens thinking, and protects human connection.

Leadership actions:

- map staff/faculty friction points before choosing tools
- pilot AI where burden is obvious: meeting notes, study plans, policy summaries, test-item drafts, curriculum mapping, inbox triage
- ask whether the tool strengthens or weakens nursing workflow
- track time saved, stress reduced, safety concerns, and relational impact
- create AI literacy expectations for faculty, staff, and students
- teach students to critique AI outputs against evidence and guidelines

Leader evaluation questions:

Does this AI tool reduce burden or add surveillance?
Does it fit the nursing process?
Does it protect human connection?
Is it explainable enough for nurses to question it?
Who is accountable if it is wrong?
What would a bedside nurse at 3 AM think?

Short workshop outline: Living and Thriving with AI

Audience: nurse leaders, educators, bedside nurses, or students.

Length: 90 minutes.

1. Opening reflection — “Where is technology helping me, and where is it draining me?”
2. Mini-teach — AI as support for nursing process, not replacement for judgment.
3. Practice — use AI to reduce a real cognitive burden: summarize, checklist, plan, or reflection.
4. Critique — identify what the AI missed, assumed, or overclaimed.
5. Systems lens — discuss workflow, equity, staffing, relationships, and safety.
6. Wellness design — create a personal AI hygiene plan: focus blocks, breaks, shutdown ritual, ergonomic adjustment.
7. Stewardship close — one boundary, one habit, one experiment for the next week.

Workshop output:

```
One AI use that reduces burden
One AI use I will avoid
One relationship I will protect
One body/attention habit I will practice
One question I will ask before trusting an AI output
```

Written guideline language

Use AI to support learning, organization, reflection, drafting, and evidence awareness. Do not use AI as a substitute for clinical judgment, patient assessment, licensed decision-making, institutional policy, or relational presence. Nurses remain accountable for verifying outputs, protecting privacy, and stewarding technology toward dignity, safety, equity, and human connection.

Google Workspace, HIPAA Boundaries, and Safe AI Use

Google Workspace matters because it is not just email. In eligible paid organizational plans, Workspace can become a centrally governed suite — Gmail, Drive, Docs, Sheets, Slides, Meet, Calendar, Chat, and related core services — under an institutional HIPAA Business Associate Agreement (BAA) when properly configured and approved.

This does **not** mean every Google account, every Google product, or every AI tool is safe for PHI.

Core rule:

Workspace may be part of a governed environment. Personal Gmail and generic external AI are not automatically governed environments.

Workspace vs personal Gmail

Environment	Practical posture
Personal Gmail / consumer Drive	No PHI, no patient stories, no employer confidential data, no proprietary organizational information.
Workspace without accepted BAA / admin governance	Treat as non-PHI for healthcare use.
Eligible paid Workspace with accepted BAA and configured controls	Potentially appropriate for PHI only when organizational privacy/compliance approves the workflow.
Gemini or AI features inside approved Workspace	Potentially usable only if included in the covered services list, configured by admin, and explicitly approved.
External generic AI tools	No PHI unless formally vetted, contracted, governed, and approved for PHI.

Why Drive, Docs, Sheets, Slides, and Meet matter

Nursing AI work often touches more than email:

- Drive stores policies, curricula, QI reports, teaching materials, meeting notes, and datasets.
- Docs is where drafts, case studies, policy language, and feedback are created.
- Sheets may hold schedules, de-identified tracking logs, or evaluation data.
- Slides carries education and leadership materials.
- Meet and Chat may include discussion, transcripts, summaries, or recordings.

If these tools are inside a governed Workspace, they can inherit central controls:

- identity and account management
- sharing restrictions
- audit logs
- retention rules
- external sharing policies
- access review

- device and download controls

That is very different from scattered personal accounts and stand-alone apps.

First-things-first PHI rule

Unless privacy/compliance has explicitly approved the tool and workflow:

PHI stays out of Hermes, personal Gmail, personal Drive, consumer AI tools, personal chatbots, and unapproved external AI.

Use instead:

- fictional cases
- de-identified teaching examples
- public policies or articles
- synthetic data
- abstracted scenarios with no unique identifiers

Safe example: nurse educator case study

Green path — non-PHI / fictional

1. Create a fictional case in an approved course planning document.
2. Ask AI to draft learning objectives, debrief questions, and assessment prompts.
3. Review for accuracy, bias, safety, and scope.
4. Store the final teaching version in the approved course folder.

Yellow / Red path — real clinical material

Stop before using external AI if the source includes:

- patient name, initials, MRN, DOB, precise dates, room numbers, rare conditions, unique story details, photos, voice, or family identifiers
- identifiable staff/student data
- employer confidential information
- operational details not approved for external sharing

Escalate to organizational privacy, security, legal, or governance review.

One-page message for nurses and faculty

Unless our privacy/compliance team has explicitly said this AI is approved for PHI and it sits in a HIPAA-covered, BAA-backed, configured environment, we use AI only with de-identified, fictional, public, or non-sensitive information.

Do / Don't examples

Scenario	Safer action	Avoid
Drafting a student case	Use a fictional composite case	Paste a real patient story into external AI
Summarizing a policy	Use public or internal-approved policy text in approved tools	Upload proprietary policy into personal AI
Creating a teaching quiz	Use fictional details and educator review	Use identifiable patient details from practice
Writing an email	Draft non-sensitive language	Include PHI in personal Gmail or generic AI
Reviewing data	Use de-identified or synthetic data unless governed	Upload spreadsheets with identifiers to unapproved AI

Workspace governance checklist

Before using Google Workspace or AI for sensitive healthcare work, confirm:

- ☐ Is this a paid, HIPAA-eligible Workspace plan?
- ☐ Has the organization accepted Google's BAA?
- ☐ Are the specific services being used covered and enabled appropriately?
- ☐ Are sharing, download, retention, and audit policies configured?
- ☐ Has privacy/compliance approved the AI feature or workflow?
- ☐ Is minimum necessary access enforced?
- ☐ Are staff trained on what cannot be pasted into AI?
- ☐ Is there a process to report accidental disclosure or unsafe use?

Hermes posture

For this course, Hermes remains a **no-PHI personal and professional learning environment** unless an institution deliberately deploys it under formal governance.

Do not describe Hermes, Google Workspace, or any AI system as HIPAA-compliant just because it is paid, enterprise, encrypted, or self-hosted. Compliance requires:

- correct environment
- access control
- auditability
- encryption
- risk analysis
- written policies
- workforce training
- vendor review
- BAA where required
- explicit organizational approval

Stewardship line:

BAA-backed tools can create a safer room. Nurses still need rules for what enters the room, who can see it, and what AI is allowed to do with it.

Advanced Growth and Sovereign Systems Pathway

The foundational class prepares participants to work responsibly with individual AI assistants and agents. As their responsibilities and projects expand, the **Advanced Growth Map** introduces the skills required to supervise an increasingly complex AI ecosystem.

Participants first learn to define agent roles, permissions, task boundaries, handoffs, escalation pathways, evaluation criteria, and human accountability. When one agent is no longer sufficient, the advanced pathway may include **OpenClaw integration** as an orchestration layer for coordinating multiple specialized agents across separate workspaces, functions, and communication channels. OpenClaw supports multi-agent routing with isolated agent workspaces, state, and session histories through a shared gateway.

Reference: <https://docs.openclaw.ai/concepts/multi-agent>

The progression is:

AI user → AI supervisor → Agent manager → Multi-agent ecosystem steward

OpenClaw integration is introduced only when multiple agents create genuine operational value. Complexity is never added for appearance, novelty, or unnecessary automation.

Every multi-agent system must retain:

- Clearly designated human ownership
- Least-privilege permissions
- Defined agent responsibilities
- Approval gates for consequential actions
- Complete logging and auditability
- Failure detection and escalation procedures
- Data-minimization and confidentiality controls
- Cost, energy, and environmental monitoring
- A manual fallback and safe shutdown process

OpenClaw or any comparable orchestration platform must not be connected to sensitive healthcare systems, patient information, clinical workflows, or organizational infrastructure without formal security review, technical validation, and institutional approval.

Separate Advanced Module: Sovereign On-Premises AI Systems

Designing and implementing a **sovereign, on-premises AI system** for a hospital unit, clinic, public-health program, or community organization is a separate advanced module and implementation pathway.

The foundational masterclass does not by itself prepare or authorize participants to deploy production clinical AI infrastructure. It provides the essential conceptual foundation needed to participate intelligently in such work, including:

- AI and agent literacy
- Human-in-the-loop workflow design
- Critical and systems thinking
- Source and output evaluation
- Privacy-aware data practices
- Role and permission design
- AI governance fundamentals

- Risk identification
- Workflow mapping
- Ethical and environmental stewardship

The sovereign-systems module builds on those foundations and addresses:

- Local and private model deployment
- Secure infrastructure architecture
- Identity and access management
- Network segmentation
- Protected health information controls
- Data residency and retention
- Encryption and key management
- Retrieval-augmented knowledge systems
- Model and agent evaluation
- Audit trails and observability
- Cybersecurity and supply-chain review
- Backup, recovery, and business continuity
- Clinical validation and change control
- Interoperability with approved organizational systems
- Incident response and safe shutdown
- Regulatory, legal, privacy, and institutional governance
- Lifecycle cost and environmental impact

Such systems require collaboration among clinical leadership, nursing, informatics, information security, privacy, legal counsel, compliance, data governance, biomedical or technical teams, and the communities affected by their use.

The class teaches participants how to begin working with AI responsibly. The Advanced Growth Map teaches them how to steward multiple agents. The Sovereign Systems Module prepares qualified teams to explore institutionally governed, locally controlled AI infrastructure.

First 10 Prompts

1. Daily Chief of Staff

Act as my nurse-centered Chief of Staff. Help me identify what matters today, what can wait, what needs follow-up, and what action protects my wellbeing.

2. Weekly Review

Run my weekly review across personal, professional, and community spheres. Help me notice patterns, risks, unfinished commitments, and one grounded next action.

3. Learning Pathway

Act as my Learning OS. Build a learning pathway for [topic] with concepts, practice, self-test, evidence check, and human-expert review points.

4. Career Map

Act as my nurse-centered career strategist. Map possible paths from my current experience toward roles that match my values, strengths, constraints, and energy.

5. Research Transformation

Use the HERMES Transformation Protocol on this article/transcript/link. Hear and Harvest, Evaluate, Reframe for Nursing, Map the System, Enable, and Steward.

6. Systems Map

Help me map this situation as a system. Identify stakeholders, incentives, feedback loops, friction points, risks, and leverage points.

7. EDENA Review

Apply the EDENA Stewardship Lens to this recommendation. Evaluate Equity and Ethics, Dignity and Data, Environment and Externalities, Nursing Relevance and Nurse Wellbeing, and Agency and Action.

8. Entrepreneurial Idea Review

Review this nurse-led idea. Be imaginative but not reckless. Identify who it helps, assumptions, risks, smallest safe pilot, evidence needed, and what I should not claim yet.

9. Dependency Check

Check whether I am over-relying on AI in this decision. What must I verify? What judgment remains mine? What human should I consult?

10. Loop Charter

Help me design a safe AI loop for [workflow]. Define trigger, goal, EDENA tier, data boundary, stop condition, human review point, and named owner.

30-Day Sustainment Plan

Weekly:

- review your Daily Brief and Weekly Review workflow
- delete or correct memory that is outdated
- check if AI is strengthening or weakening your judgment
- update one prompt or template
- identify one thing that should remain human-only

Monthly:

- review your no-PHI boundary
- review your professional goals
- apply EDENA to one major recommendation
- retire one workflow that no longer helps
- choose one safe improvement for the next month

Completion Statement

I have built a nurse-centered AI operating system for personal, professional, and community use, governed by human agency, evidence-awareness, safety boundaries, and nursing stewardship.

The Closing Charge: Be a Nightingale for Your Time

Florence Nightingale did not choose between the heart and the evidence — she held both, and that is why the wards changed. Your generation inherits her charge with new instruments. Data science and intelligent agents can now carry the counting, the sorting, the drafting, and the remembering, so that you can carry what only a nurse can: presence, judgment, and care for the human being in front of you.

Carry the lamp. Keep the ledger. Bring your light on the human condition — and let the system do the rest.

Agents propose. Humans judge. Nurses steward.

Nurse AI Side-Gig Starter Kit — First Doorway into NIN/NAIO

Purpose: help nurses explore one safe, no-PHI AI-enabled side project before stepping into the deeper NIN/NAIO community, Founding Steward Cohort, Builder Academy, or EDENA governance pathway.

This is not a generic digital-product kit. It is a nurse-centered validation system for templates, prompt packs, workshops, content systems, study aids, dashboards, and community resources that stay outside patient-specific care.

What is included

- EDENA no-PHI guardrails.
- Side-Gig Offer Canvas.
- 30-Day Validation Sprint.
- Prompt Pack.
- 30-Day Trust-Building Content Calendar.
- Landing Page Copy.
- Outreach Conversation Script.

Safe positioning

Build your first safe, no-PHI nurse-AI project in 30 days — with prompts, templates, guardrails, and a guided operating system.

Pathway

Starter Kit → Lamp Huddle → NIN community → Founding Steward Cohort → Builder Academy → EDENA governance pathway.

Boundary

No PHI. No patient-specific clinical decisions. No employer secrets. No certification claim. Human judgment first.

First action

Open `03-Community-Entrepreneurship/Nurse-AI-Side-Gig-Starter-Kit/Side-Gig-Offer-Canvas.md`, choose one buyer, write one safe promise, and talk to 10 nurses before building a large product library.

nurse-ai-os.org · No PHI, ever · AI drafts, humans judge, nurses steward · This PDF is generated from assets/nurse-ai-os-workbook.md — the web copy is canonical.