

AI Operating System with the Heart of a Nurse

Starter Kit + Course Workbook

Course promise: Build a nurse-centered AI operating system for personal, professional, and community / entrepreneurial use - governed by human agency, evidence-awareness, systems thinking, ethical stewardship, safety, and dignity.

Default posture: Hermes becomes your nurse-centered Chief of Staff, Learning OS, strategic advisor, and stewardship partner.

Core line:

> Hermes supports. Humans judge. Nurses steward.

Read This First: Boundaries

This course is for personal productivity, learning, research, writing, reflection, career planning, community building, and non-PHI professional work.

It is not for:

- patient-specific clinical decisions
- medication, diagnosis, or treatment recommendations
- EHR screenshots
- patient names, MRNs, chart excerpts, or identifiers
- confidential employer information without authorization
- replacing licensed judgment or institutional policy

If the work could affect patient care, expose PHI, or create clinical accountability, pause and route through appropriate human and institutional governance.

The Nurse-Centered AI OS Stack

Heart of Nursing

v

Human Agency

v

Evidence Awareness

v

Systems Thinking

v

Ethical Governance

v
Actionable Workflows
v
Safe Execution
v
Reflective Learning

A nurse-centered AI OS should be:

- nurse-centered
- human-agency preserving
- evidence-aware
- systems-thinking oriented
- ethically governed
- actionable
- safety-first
- dignity-protecting
- entrepreneurially imaginative but not reckless

Course Map

Week 1 - Heart Before Tools

Goal: Define the values and boundaries that will guide your AI OS.

Build

- MY-AI-VALUES.md
- NO-PHI-BOUNDARY.md
- HUMAN-AGENCY-RULES.md

Prompt

Help me write a personal AI values statement as a nurse.

It should reflect:

- human judgment
- privacy
- dignity
- evidence-awareness
- critical thinking
- safety
- wellbeing
- service

- professional boundaries

Use clear language I can paste into my AI OS setup.

Week 2 - Set Up Hermes Desktop Safely

Goal: Install Hermes Desktop and create a nurse-centered AI workspace.

Build

- Hermes Desktop installed
- nurse-centered profile or working folder
- starter markdown architecture
- no-PHI safety boundary

Starter folder

My-Nurse-AI-OS/

00-Start-Here/

01-Personal/

02-Professional/

03-Community-Entrepreneurship/

04-Governance/

05-Templates/

Prompt

Help me set up my Nurse-Centered AI Operating System folder.

Create a simple structure for personal, professional, and community / entrepreneurial support.

Include safety boundaries, human-agency rules, and weekly review templates.

Hermes Desktop Setup Walkthrough

Use the setup tutorial as a practical orientation, but translate it through the nurse-centered safety lens.

1. Install and launch Hermes

For macOS, the easiest beginner path is the official Hermes Desktop installer.

- Go to the official Hermes Agent site: <https://hermes-agent.nousresearch.com/>
- Download the macOS Desktop installer from the official site or official Nous Research GitHub release path.
- Avoid third-party mirrors and unofficial download sites.
- Open the downloaded .dmg file.
- Drag Hermes into Applications when prompted.
- Eject the DMG, then launch Hermes from Applications or Spotlight.
- If Gatekeeper warns that the app came from the internet, confirm it came from the official Hermes/Nous source before opening.

The Desktop app and CLI share the same Hermes core: config, provider settings, sessions, skills, memory, and profiles. A setup completed in Desktop also prepares the underlying Hermes CLI for later advanced use.

If you are comfortable with Terminal, you can verify the terminal UI after setup:

```
hermes --tui
```

If the command does not work, use the official docs path first: run Hermes Desktop onboarding, or use the official hermes setup flow to repair provider/config issues.

Beginner note: current platform support and installer details can change. Always use the official Hermes Agent docs as the source of truth.

After launch:

- Open a new session.
- Notice the main areas: sessions, pinned messages, artifacts, skills, tools, messaging, settings, and provider/model settings.

2. Configure the workspace

Set the workspace to a folder you intentionally created for this course, such as:

```
My-Nurse-AI-OS/
```

Do not point Hermes at folders that contain patient information, confidential employer files, tax records, passwords, or sensitive family documents unless you fully understand and approve the access.

3. Choose safety controls

Start with manual approval, not full autonomy.

For nurses, the recommended beginner posture is:

Manual approval first.

Small safe tasks first.

No PHI.

No clinical decisions.

No background automation until you understand the risk.

Smart or autonomous approval settings should be treated as advanced features. More access means more governance.

4. Understand memory

Hermes can remember durable preferences and context. This is powerful, but nurses should treat memory as a governed space.

Use memory for:

- writing preferences
- learning goals
- non-PHI professional interests
- course progress
- recurring workflow preferences

Do not store:

- patient identifiers
- confidential employer details
- secrets, passwords, or API keys
- sensitive family information
- anything you would not want resurfacing later

5. Configure models and providers

Hermes can connect to providers such as OpenRouter, Google Gemini, GitHub Copilot, Anthropic, and others depending on current Hermes support and your account access.

Some providers may list free or low-cost models. Treat "free" as temporary and subject to change. Always check current provider terms, privacy policy, rate limits, and cost before using a model for serious work.

Never paste API keys into public notes, screenshots, course forums, shared documents, or GitHub repos. If a key appears in a file, revoke it and create a new one.

6. Test with a safe task

Start with a low-risk task such as:

Read my Nurse-AI-OS folder and summarize the starter files you see. Do not access other folders. Do not modify anything yet.

Then try a bounded file task:

Create a new folder called Practice inside My-Nurse-AI-OS and write a short README explaining that this is a safe test folder.

7. Use sub-agents responsibly

Sub-agents can research different sides of a question in parallel while Hermes supervises and synthesizes the result.

Safe beginner example:

Fire two sub-agents for a non-clinical learning task.

Agent 1: summarize the benefits of using AI for nurse professional development.

Agent 2: summarize the risks and limitations.

Then synthesize both sides and tell me what I should verify before acting.

No PHI. No clinical decisions.

Avoid using sub-agents for patient-specific care, confidential employer work, or decisions with financial, legal, clinical, or employment consequences without human expert review.

8. Voice, messaging, and connectors

Voice, Telegram, Discord, Slack, and other messaging integrations can make Hermes more useful, but each connector creates a new doorway.

Before connecting a channel, ask:

- What data can pass through this channel?
- Who else can see it?
- Can messages be forwarded, stored, or searched?
- What should never be sent here?
- How do I disconnect it?

Connector rule:

> Do not connect Hermes to sensitive healthcare systems, employer systems, patient data, or confidential channels without formal approval.

macOS Day-One Checklist

Use these first tasks after installation:

1. Verify a normal chat works

Give me a 5-bullet orientation to this Nurse-AI-OS course. Do not use clinical advice. Do not ask for PHI.

2. Create a safe daily brief

Act as my nurse-centered Chief of Staff. Create a daily brief for personal priorities, professional learning, and one community or entrepreneurial action. Do not include patient information.

3. Check your workspace boundary

Tell me what workspace folder you can see. Do not read outside the course folder. Tell me if anything looks sensitive before opening it.

4. Try a balanced sub-agent research task

Use two sub-agents for a non-clinical learning question. One should summarize benefits; one should summarize risks. Synthesize the result and tell me what requires human verification.

5. Do not add advanced features yet

Do not enable computer use, messaging connectors, cron jobs, background loops, or autonomous approval until a normal chat, workspace boundary, and no-PHI workflow are working reliably.

Personal Projects and Personal Learning Profiles

Hermes profiles let you run separate agents on the same machine. Each profile has its own configuration, .env, SOUL.md, memories, sessions, skills, cron jobs, state database, gateway state, and identity.

Use profiles to avoid mixing personal learning, personal projects, and professional work.

Recommended beginner profiles:

Personal Projects
Personal Learning
Professional Non-PHI

Important boundary:

> Profiles help separate memory, configuration, skills, and sessions. They are not a substitute for data governance, access control, or clinical privacy review.

Personal Projects profile

Use this for side projects, writing, small apps, personal goals, and low-stakes experiments.

Suggested SOUL.md:

Personal Projects SOUL

You help me plan and ship small personal side projects and learning experiments.
You keep my goals realistic, help me choose next actions, and protect my time and energy.
You do not access work files, patient information, or confidential employer material.
You preserve my judgment and ask for approval before changing files or creating scheduled jobs.

Start with one recurring personal workflow:

- weekly project review
- daily idea grooming
- lightweight learning plan
- personal project next-action list

Day-one prompt:

I am in my Personal Projects profile.
Here are my current personal projects and rough goals: [PASTE LIST]
Design one simple weekly review workflow with clear input, output, manual approval, and

a reusable prompt.

Do not create a cron job yet. Do not access work or clinical files.

Personal Learning profile

Use this for books, articles, podcasts, videos, courses, research notes, and your second brain.

Suggested SOUL.md:

Personal Learning SOUL

You help me learn efficiently, turn reading into structured notes, and surface past insights when relevant.

You are evidence-aware, systems-thinking oriented, and careful about uncertainty.

You distinguish source claims, assumptions, implications, and my own reflections.

You do not turn learning notes into clinical recommendations.

Choose one knowledge home:

- a local markdown folder such as ~/Notes/Learning
- an Obsidian vault
- a course folder inside My-Nurse-AI-OS

Create an inbox pattern:

Learning/

Inbox/

Notes/

Weekly-Digests/

Skills/

Knowledge inbox prompt:

I am in my Personal Learning profile.

Process this as a learning inbox item.

Use the HERMES Transformation Protocol:

1. Hear and Harvest the source faithfully.
2. Evaluate claims, assumptions, uncertainty, and bias.
3. Reframe for nursing, wellbeing, leadership, policy, AI, or innovation.
4. Map the system: stakeholders, incentives, feedback loops, constraints, and leverage points.

5. Enable: create prompts, checklists, workflows, learning steps, or experiments.
6. Steward: name safety, privacy, equity, accountability, reversibility, and stop conditions.

Output a markdown note with:

- title and date
- 5-10 bullet summary
- key concepts and definitions
- what this changes for my thinking or practice
- open questions
- what I should verify
- related notes or tags

Do not use PHI. Do not make patient-specific recommendations.

Saving workflows as skills

When a workflow works after 2-3 uses, turn it into a skill or reusable template.

Good candidates:

- process-learning-inbox
- weekly-project-review
- research-claim-check
- career-map-review

Skill rule:

> Save repeated procedures as skills only after the workflow has proven useful. Do not create a skill for every one-off prompt.

Weekly review cron - only after the manual workflow works

Cron jobs run in fresh sessions, so prompts must be self-contained. Do not assume the cron job remembers the current chat.

Begin with a manual weekly review. If the output is consistently useful, then create a scheduled job.

Safe personal-learning cron concept:

Every Sunday at 5pm, review the Learning/Notes folder for new or updated notes from the past 7 days. Create a weekly learning digest with: what I read, top 5 ideas, 3 follow-up questions, and 1 grounded learning focus for next week. Do not access clinical or work files. Do not process PHI. Save the digest to Learning/Weekly-Digests/.

Safe personal-project cron concept:

Every Sunday at 5pm, review the Personal Projects folder and create a weekly project review with: completed progress, blocked items, highest-leverage next actions, and one thing to pause. Do not make purchases, send messages, change files outside this folder, or create new automations without approval.

Cron boundary:

> Schedule only low-stakes, non-PHI, review-oriented workflows at first. Anything involving messaging, external systems, financial actions, legal decisions, clinical work, or employer systems requires a higher EDENA tier and human approval.

Week 3 - Build Your Nurse-Centered Chief of Staff

Goal: Create a daily and weekly operating rhythm.

Daily Brief Prompt

Act as my nurse-centered Chief of Staff.

Create my daily brief using this structure:

1. What matters today
2. What needs preparation
3. What could create stress or overload
4. What I should not forget
5. What can wait
6. One action that protects my wellbeing
7. One decision that remains mine

Weekly Review Prompt

Run my weekly nurse-centered review.

Use this structure:

1. What happened this week?
2. What did I carry well?
3. What drained me?
4. What needs follow-up?
5. What pattern should I notice?

6. What relationship or responsibility needs attention?
7. What should I stop, start, or continue?
8. What is the safest next step?

Week 4 - Build Your Learning OS

Goal: Turn information into capability without outsourcing judgment.

Learning Pathway Prompt

Act as my nurse-centered Learning OS.

I want to learn: [TOPIC]

Build a learning pathway that includes:

1. What I need to understand first
2. Core concepts
3. Common misconceptions
4. Practice scenarios
5. Questions to test my understanding
6. What evidence or guidelines I should review
7. What I should ask a human expert
8. How I will know I am improving
9. What I should not use AI for in this topic

Week 5 - Build Your Professional + Community OS

Goal: Map career growth, contribution, and nurse-led innovation.

Career Map Prompt

Act as my nurse-centered career strategist.

Use my background, values, interests, constraints, and energy realities to create a career map.

Include:

1. Current strengths
2. Possible paths
3. Skills to build

4. Credentials to consider
5. People to learn from
6. Risks of each path
7. First small experiment
8. What decision remains mine

Entrepreneurial Idea Prompt

Act as my nurse-centered entrepreneurship advisor.

I have this idea: [IDEA]

Help me explore it without becoming reckless.

Use this structure:

1. Who might this help?
2. What problem does it address?
3. What evidence do we have?
4. What assumptions are we making?
5. What could go wrong?
6. What ethical or equity issues matter?
7. What is the smallest safe test?
8. What should not be promised yet?
9. What human review is needed?
10. What is the next grounded action?

Week 6 - Govern, Review, and Sustain Your AI OS

Goal: Build long-term trust, boundaries, and review habits.

HERMES Transformation Protocol

Use this for transcripts, links, articles, lectures, podcasts, ideas, or research.

H - Hear and Harvest

Faithfully understand the source before remixing it.

E - Evaluate

Assess evidence, assumptions, uncertainty, and bias.

R - Reframe for Nursing

Translate into bedside, education, leadership, AI, policy, wellbeing, or innovation.

M - Map the System

Identify stakeholders, incentives, feedback loops, constraints, and leverage points.

E - Enable

Create tools, prompts, checklists, workflows, pilots, or learning plans.

S - Steward

Check safety, privacy, equity, accountability, reversibility, and stop conditions.

EDENA Stewardship Lens

Use this for major recommendations.

E - Equity and Ethics

Who benefits? Who could be excluded or harmed?

D - Dignity and Data

What data is involved? Is privacy protected? Is dignity preserved?

E - Environment and Externalities

What effects appear outside the immediate task?

N - Nursing Relevance and Nurse Wellbeing

Does this support nursing judgment, workflow, and wellbeing?

A - Agency and Action

Who remains accountable? What action is safe, reversible, and human-led?

Governed Loop Charter

Every loop needs a leash, a log, a limit, and a human owner.

Loop Name:
Purpose:
Trigger:
Goal:
EDENA Tier:
Data Used:
PHI Risk:
Tools/Connectors:
Evaluation Method:
Stop Condition:
Maximum Runtime:
Maximum Cost:
Human Review Point:
Escalation Condition:
Audit Log Location:
Rollback Plan:
Named Human Owner:

Advanced Growth and Sovereign Systems Pathway

The foundational class prepares participants to work responsibly with individual AI assistants and agents. As their responsibilities and projects expand, the Advanced Growth Map introduces the skills required to supervise an increasingly complex AI ecosystem.

Participants first learn to define agent roles, permissions, task boundaries, handoffs, escalation pathways, evaluation criteria, and human accountability. When one agent is no longer sufficient, the advanced pathway may include OpenClaw integration as an orchestration layer for coordinating multiple specialized agents across separate workspaces, functions, and communication channels. OpenClaw supports multi-agent routing with isolated agent workspaces, state, and session histories through a shared gateway.

Reference: <https://docs.openclaw.ai/concepts/multi-agent>

The progression is:

AI user -> AI supervisor -> Agent manager -> Multi-agent ecosystem steward

OpenClaw integration is introduced only when multiple agents create genuine operational value. Complexity is never added for appearance, novelty, or unnecessary automation.

Every multi-agent system must retain:

- Clearly designated human ownership
- Least-privilege permissions
- Defined agent responsibilities
- Approval gates for consequential actions
- Complete logging and auditability
- Failure detection and escalation procedures
- Data-minimization and confidentiality controls
- Cost, energy, and environmental monitoring
- A manual fallback and safe shutdown process

OpenClaw or any comparable orchestration platform must not be connected to sensitive healthcare systems, patient information, clinical workflows, or organizational infrastructure without formal security review, technical validation, and institutional approval.

Separate Advanced Module: Sovereign On-Premises AI Systems

Designing and implementing a sovereign, on-premises AI system for a hospital unit, clinic, public-health program, or community organization is a separate advanced module and implementation pathway.

The foundational masterclass does not by itself prepare or authorize participants to deploy production clinical AI infrastructure. It provides the essential conceptual foundation needed to participate intelligently in such work, including:

- AI and agent literacy
- Human-in-the-loop workflow design
- Critical and systems thinking
- Source and output evaluation
- Privacy-aware data practices
- Role and permission design
- AI governance fundamentals
- Risk identification
- Workflow mapping
- Ethical and environmental stewardship

The sovereign-systems module builds on those foundations and addresses:

- Local and private model deployment
- Secure infrastructure architecture
- Identity and access management
- Network segmentation
- Protected health information controls
- Data residency and retention
- Encryption and key management
- Retrieval-augmented knowledge systems
- Model and agent evaluation
- Audit trails and observability

- Cybersecurity and supply-chain review
- Backup, recovery, and business continuity
- Clinical validation and change control
- Interoperability with approved organizational systems
- Incident response and safe shutdown
- Regulatory, legal, privacy, and institutional governance
- Lifecycle cost and environmental impact

Such systems require collaboration among clinical leadership, nursing, informatics, information security, privacy, legal counsel, compliance, data governance, biomedical or technical teams, and the communities affected by their use.

> The class teaches participants how to begin working with AI responsibly. The Advanced Growth Map teaches them how to steward multiple agents. The Sovereign Systems Module prepares qualified teams to explore institutionally governed, locally controlled AI infrastructure.

First 10 Prompts

1. Daily Chief of Staff

Act as my nurse-centered Chief of Staff. Help me identify what matters today, what can wait, what needs follow-up, and what action protects my wellbeing.

2. Weekly Review

Run my weekly review across personal, professional, and community spheres. Help me notice patterns, risks, unfinished commitments, and one grounded next action.

3. Learning Pathway

Act as my Learning OS. Build a learning pathway for [topic] with concepts, practice, self-test, evidence check, and human-expert review points.

4. Career Map

Act as my nurse-centered career strategist. Map possible paths from my current experience toward roles that match my values, strengths, constraints, and energy.

5. Research Transformation

Use the HERMES Transformation Protocol on this article/transcript/link. Hear and Harvest, Evaluate, Reframe for Nursing, Map the System, Enable, and Steward.

6. Systems Map

Help me map this situation as a system. Identify stakeholders, incentives, feedback loops, friction points, risks, and leverage points.

7. EDENA Review

Apply the EDENA Stewardship Lens to this recommendation. Evaluate Equity and Ethics, Dignity and Data, Environment and Externalities, Nursing Relevance and Nurse Wellbeing, and Agency and Action.

8. Entrepreneurial Idea Review

Review this nurse-led idea. Be imaginative but not reckless. Identify who it helps, assumptions, risks, smallest safe pilot, evidence needed, and what I should not claim yet.

9. Dependency Check

Check whether I am over-relying on AI in this decision. What must I verify? What judgment remains mine? What human should I consult?

10. Loop Charter

Help me design a safe AI loop for [workflow]. Define trigger, goal, EDENA tier, data boundary, stop condition, human review point, and named owner.

30-Day Sustainment Plan

Weekly:

- review your Daily Brief and Weekly Review workflow
- delete or correct memory that is outdated
- check if AI is strengthening or weakening your judgment
- update one prompt or template
- identify one thing that should remain human-only

Monthly:

- review your no-PHI boundary
- review your professional goals
- apply EDENA to one major recommendation
- retire one workflow that no longer helps
- choose one safe improvement for the next month

Completion Statement

I have built a nurse-centered AI operating system for personal, professional, and community use, governed by human agency, evidence-awareness, safety boundaries, and nursing stewardship.