

The Plain-Language Safety Rules

Nurse AI OS · Safety & Governance. You memorized the rights of medication administration. These are shorter — and they protect your patients, your license, and you, at every level of AI use.

Agents propose. Humans judge. Nurses steward.

Free for nurses and nursing students · Founding Year 2026 · nurse-ai-os.org

The five rules

The one-minute test before entering anything sensitive: if you cannot name the approved environment, the data type, the AI feature, and the accountable human owner — don't put it in the tool.

1 · No patient data. Ever.

No names, charts, screenshots, room numbers, record numbers, dates, or stories that could identify a patient — in any AI tool, at any level of this program. If you're unsure whether something counts: it counts. Leave it out.

2 · AI drafts. You decide.

Every AI output is a draft until a human reviews it. Nothing gets sent, published, installed, or acted on without your judgment. Human gates are checkpoints, not suggestions.

3 · Your license is not outsourceable.

No AI here supports clinical decisions, assessments, or patient-specific care questions. The system carries your busywork — never your judgment, never your accountability.

4 · Employer systems stay at work.

No EHR screens, no confidential policies, no identifiable staff or student data, no workplace secrets. Public, fictional, de-identified, or personal material only — unless your organization's compliance team has explicitly approved a tool and workflow.

5 • When in doubt, ask a human.

Bring gray areas to the monthly Lamp Huddle or email great.ai.nurses@gmail.com. A question asked early beats an incident reported late.

These rules have a name: EDENA

The five rules above are the plain-language version of **EDENA — Ethical Design & Enablement for Nursing Augmentation** — the governance framework enshrined in every layer of Nurse AI OS. AI may carry workflows; humans remain responsible for judgment. Every action you ever automate gets classified before it runs. Tiers measure **risk, not privilege**:

- **GREEN** — low risk, may proceed
- **YELLOW** — conditional, you approve first
- **RED** — stop; human authorization mandatory, always

(A fourth tier, Orange, covers advanced integrations like APIs and automations. It is deferred — treated as Red until a governed review defines its scope — and you won't meet it in your first month.)

You don't need to memorize this. The 30-Day Roadmap tags every task with its tier, and you stay in Green and Yellow for your whole first month. By the time you meet EDENA formally, it won't feel like a framework — it will feel like how you already think: assess the risk before you act, and keep a human hand on anything that matters.

Before installing any new routine — the seven questions

Purpose · Data involved · PHI? (if yes: stop) · Autonomy level · Human approval point · What failure looks like · Go or no-go. If you can't answer all seven, the routine isn't ready.

Carry the lamp. Keep the ledger.

© NAIOS Institute · Not medical advice · Not for PHI · Not a replacement for licensed judgment · nurse-ai-os.org/setup.html#safety

nurse-ai-os.org · No PHI, ever · AI drafts, humans judge, nurses steward · This PDF is generated from `assets/safety-rules-edena.md` — the web copy is canonical.